

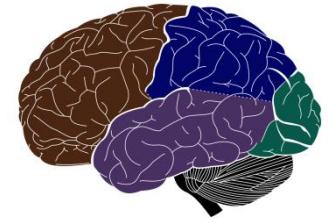
# Neural networks (NN)

Ruxandra Stoean

# Additional bibliography

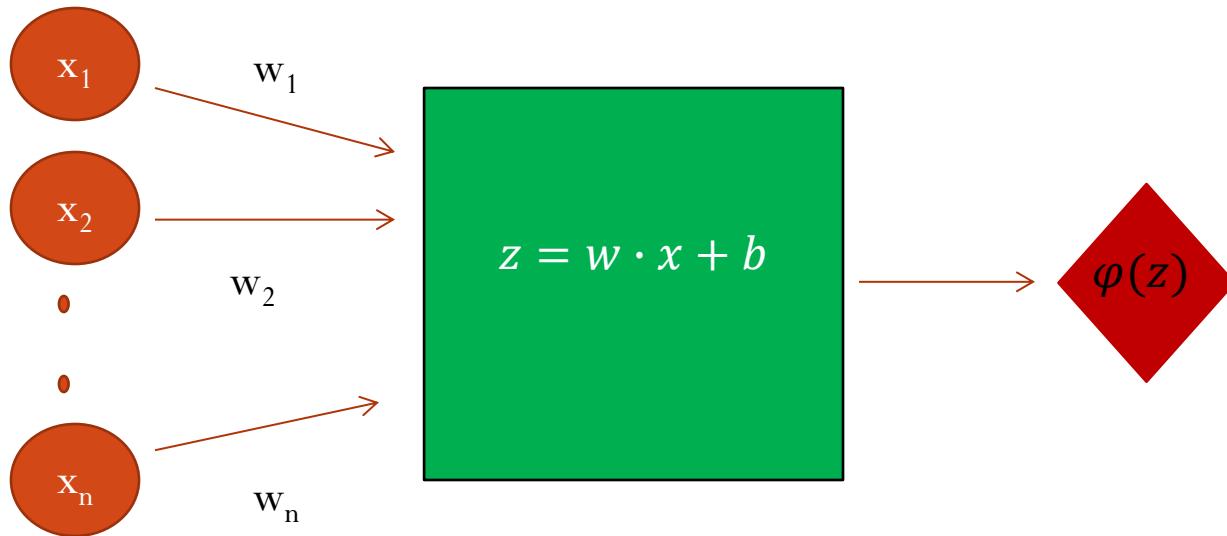
- Simon O. Haykin , Neural Networks and Learning Machines (3rd Edition), Prentice Hall, 2008
- Charu C. Aggarwal, Neural Networks and Deep Learning: A Textbook, Springer, 2018
- Christopher M. Bishop, Neural Networks for Pattern Recognition, Oxford University Press, 1996

# Modelling the human brain



- Artificial neural networks simulate the mode of neural interaction within the human brain, directed towards learning:
  - Base units – **neurons**
  - Connections between them – **synapses**
- The intensities (weights) of the synapses determine the performance of learning.

# Artificial neuron McCulloch-Pitts



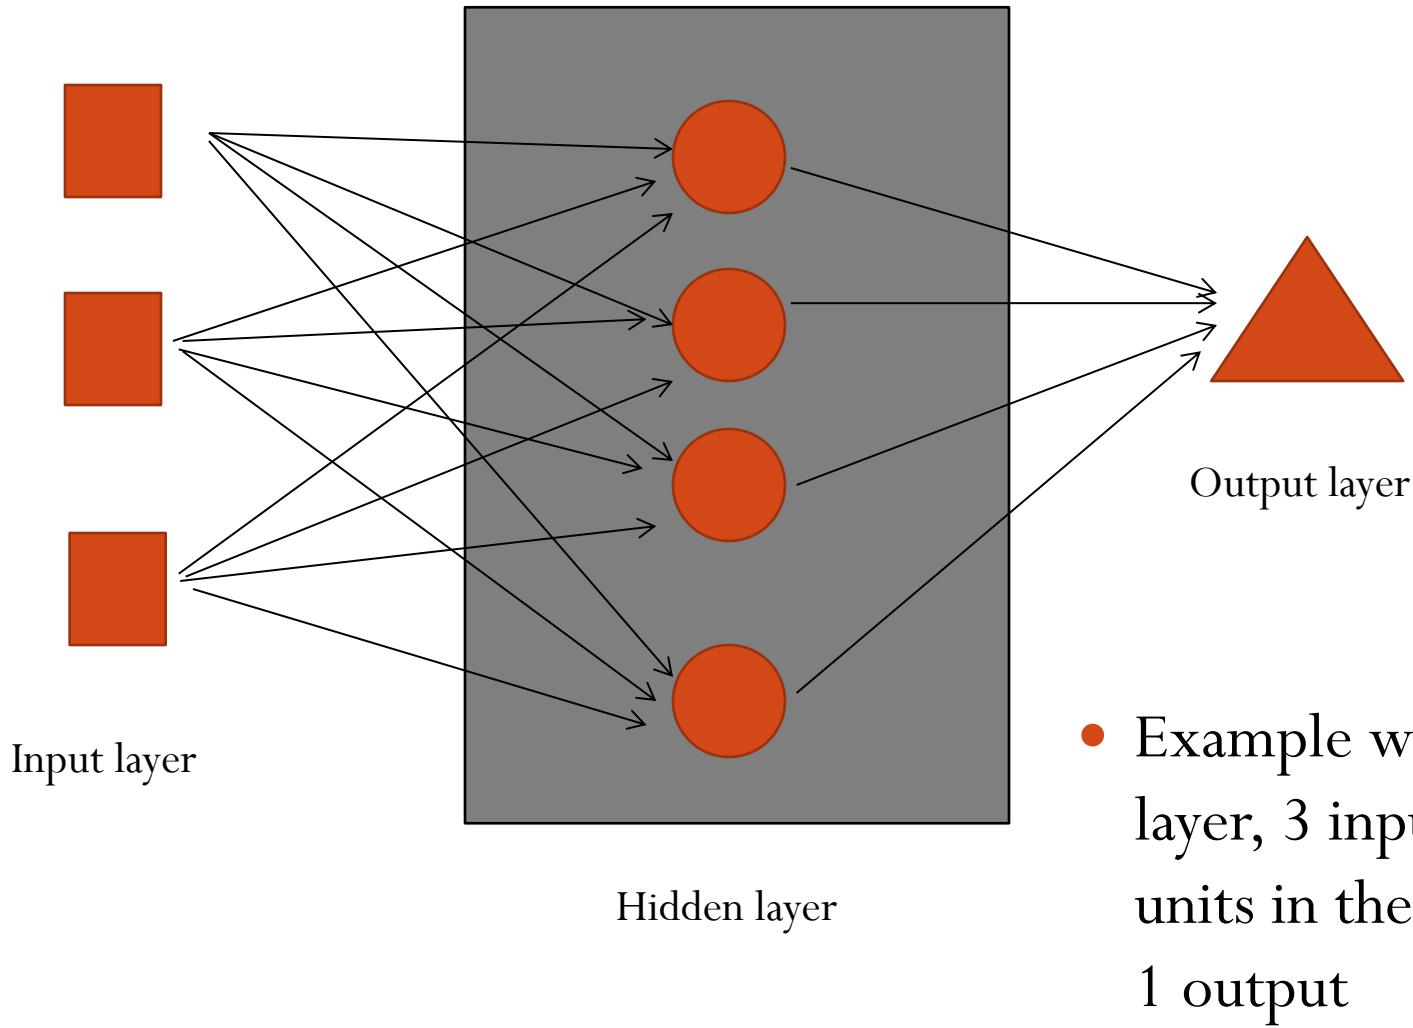
- $x$  – input vector
- $w$  – synaptic weights
- $b$  – bias

- $z$  – linear combination unit
- $\varphi$  – activation function of the neuron
- Output  $\varphi(z)$

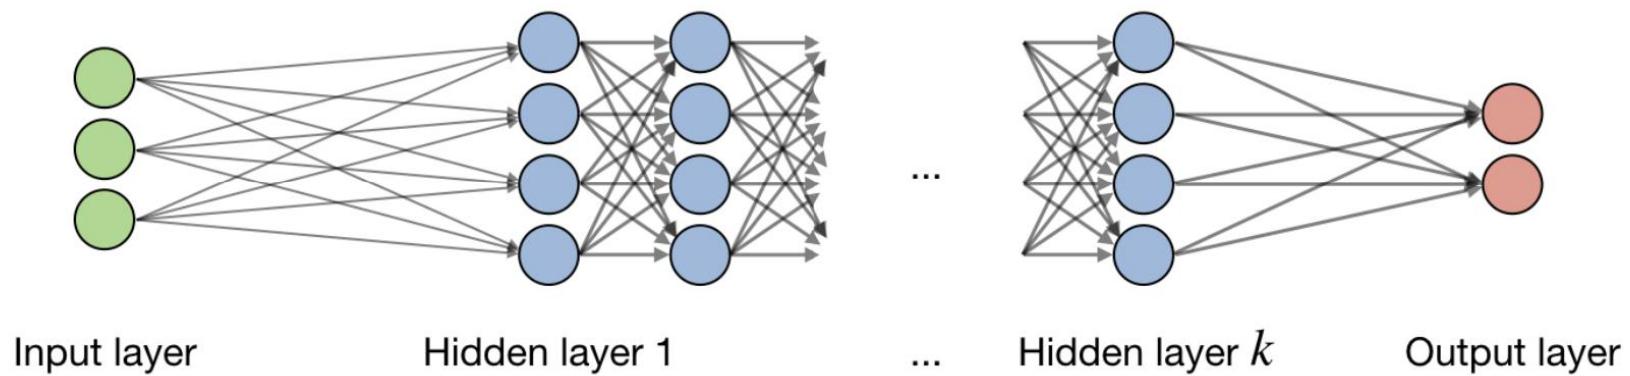
# Structure of a neural network

- Input neurons (units) – input layer (input data attributes)
- Hidden neurons in the black-box of learning (inside data components to be learnt)
  - In one or more hidden layers
  - The output of a layer becomes the input for the next layer
- Output neurons – output layer (network output - classes/response)
- The (supervised) learning starts from the problem data and optimizes the weights of the neurons on the base of the difference between the predicted and the real response.

# Simple NN architecture



# General NN architecture



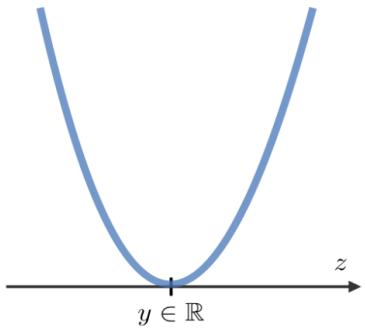
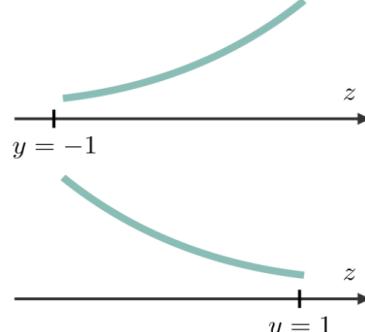
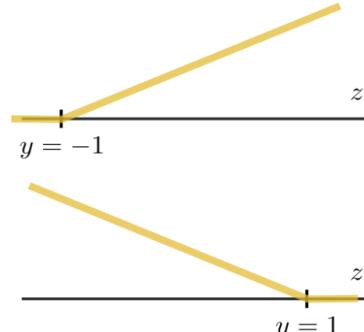
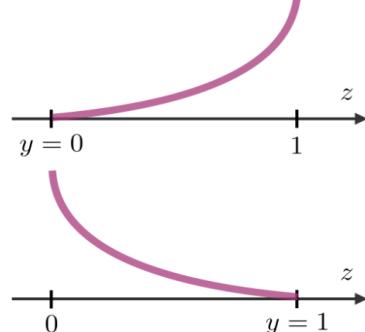
<https://stanford.edu/~shervine/teaching/cs-229>

# NN Flow

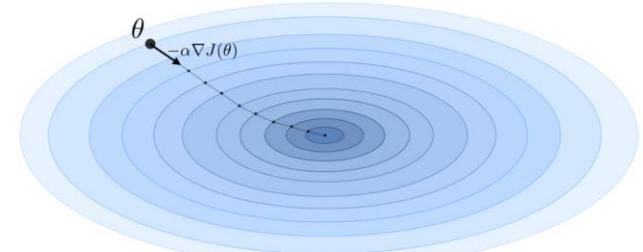
- Batches of input data are given in the input layer
  - Batch size – amount of data passed in one forward pass iteration
- Neurons are multiplied by their weights
  - The product is further transformed by the activation function
  - The new collection of neurons are fed to the next layer
- The output of the last layer – the prediction – is compared to the output -> the loss -> to be minimized
  - Weights are initially randomly initialized
  - But, with the loss, the gradient of the cost function is calculated for the weights
  - Weights are updated after each iteration
- All data batches pass – one epoch ends
- The process is repeated for a number of epochs

# Loss function

- The difference between the predicted value of the model  $z$  and the corresponding data real output  $y$
- Cost function – sum of loss over all training data (or batch)
- NN use (binary, categorical) cross-entropy loss since it allows probability estimates; MSE for regression

Least squared error	Logistic loss	Hinge loss	Cross-entropy
$\frac{1}{2}(y - z)^2$	$\log(1 + \exp(-yz))$	$\max(0, 1 - yz)$	$-(y \log(z) + (1 - y) \log(1 - z))$
			
Linear regression	Logistic regression	SVM	Neural Network

# Backpropagation



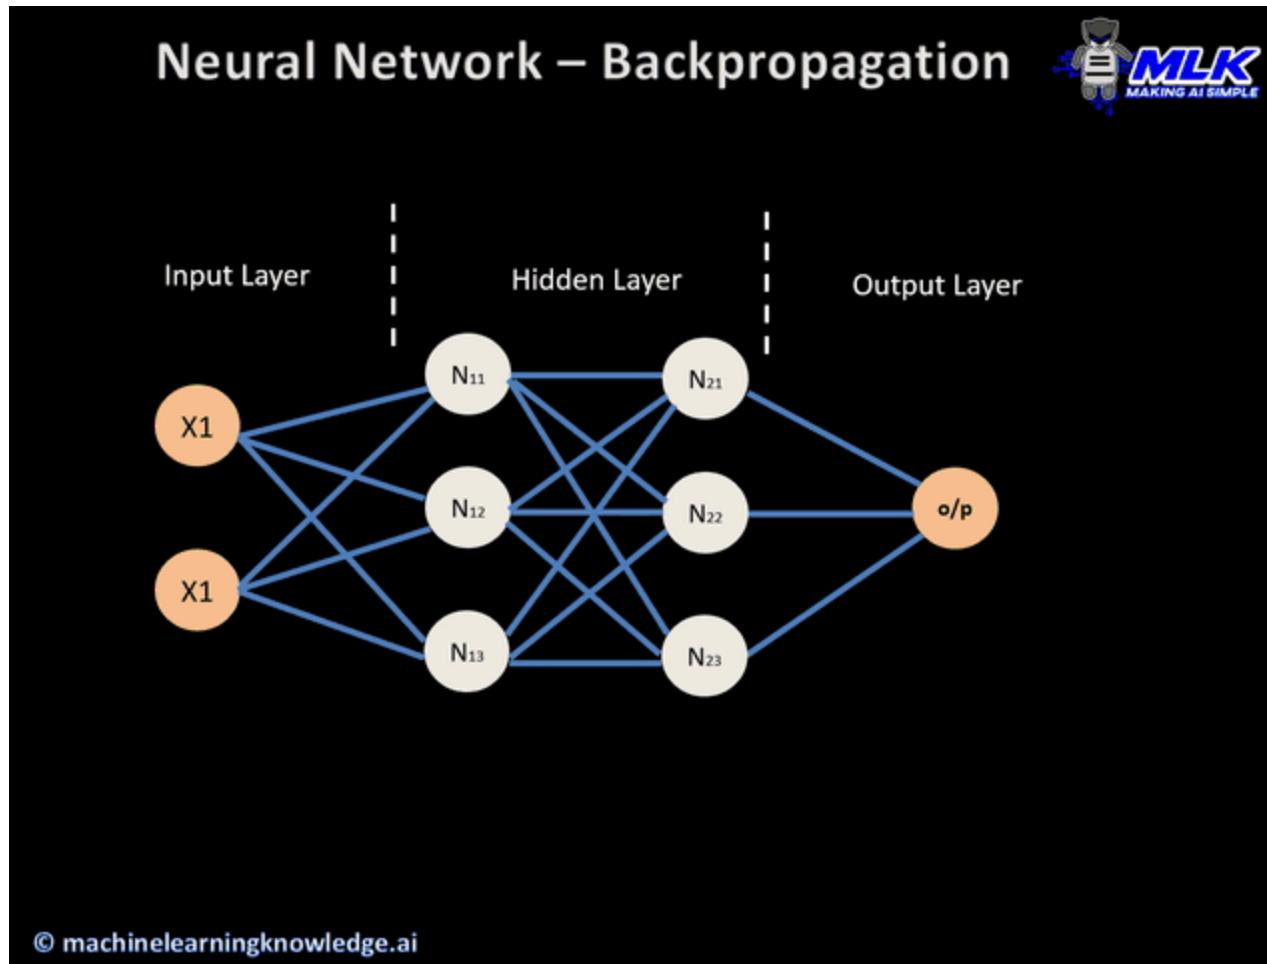
<https://stanford.edu/~shervine/teaching/cs-229>

- Backpropagation
  - The process of propagating the error backward in order to update the weights
  - It computes the gradient of the cost function with respect to the model weights through the chain rule from calculus
- Gradient descent
  - Optimization algorithm to find the best weight values based on the gradient
  - Descends down the cost function, by the learning rate, to the minimum
  - Examples: Stochastic gradient descent, Adam, AdaGrad, RMSProp

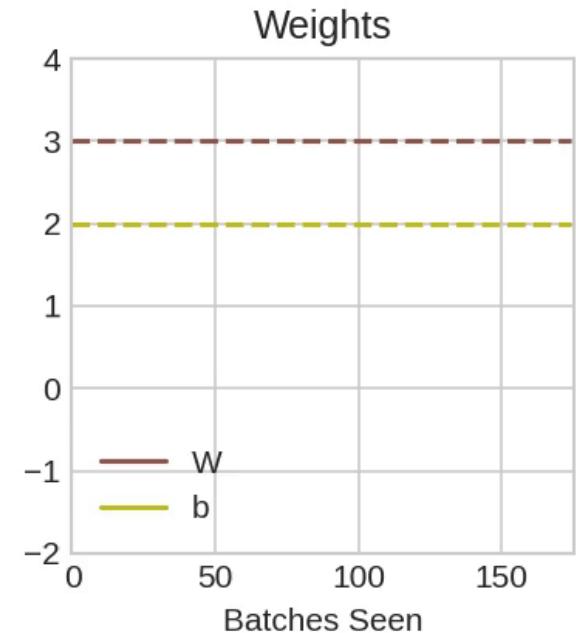
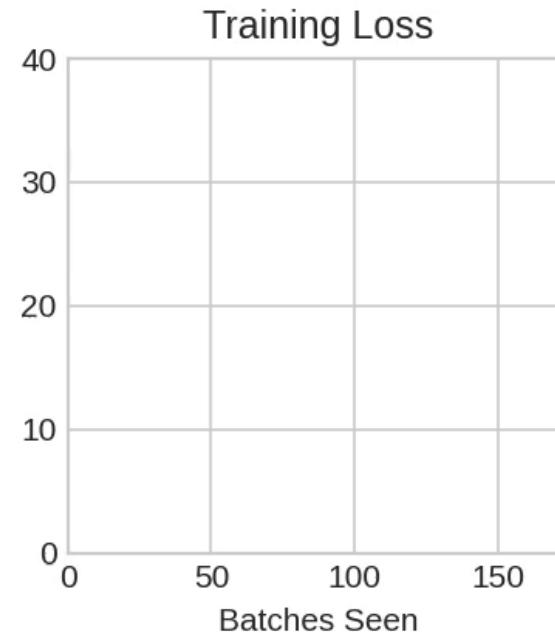
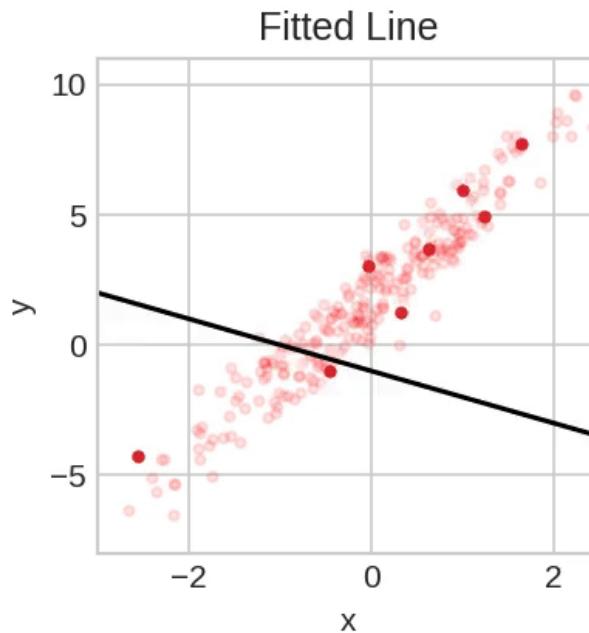
# Learning rate

- The rate of updating the weights
  - Constant
  - Small – converges slowly and gets stuck in local minima
  - Large – unstable
- Better let it decrease gradually by time step
- Or even better – adaptive (e.g. Adam), depending on the decrease of the training loss

# Illustration



# Training with SGD – batches in bold



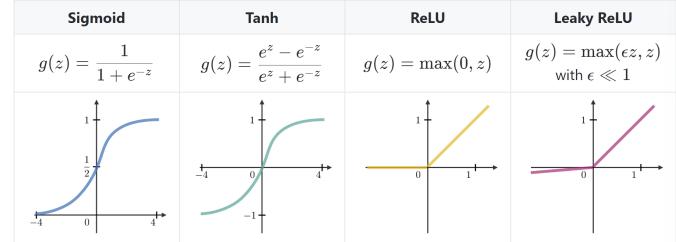
<https://www.kaggle.com/learn>

# Activation function

- It induces non-linearity to the computed unit  $z$ .

Sigmoid	Tanh	ReLU	Leaky ReLU
$g(z) = \frac{1}{1 + e^{-z}}$	$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z) \text{ with } \epsilon \ll 1$

# Choice of activation

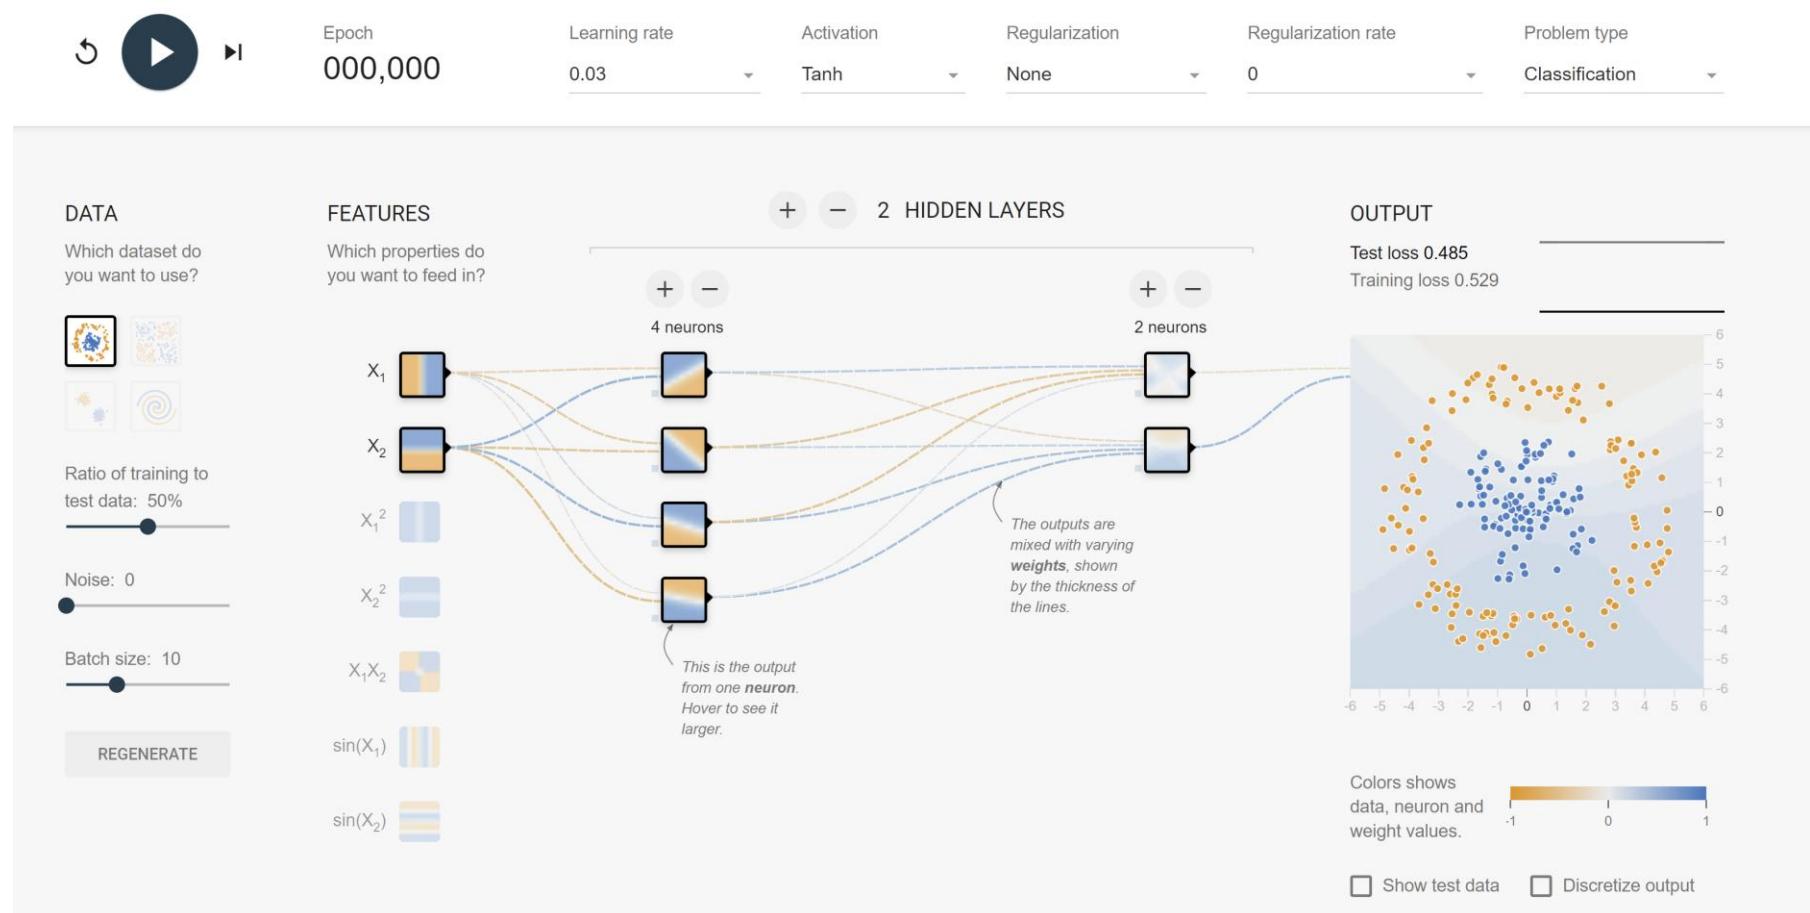


- **Sigmoid (logistic)**: smooth gradient; normalized output (between 0 and 1); clear prediction
  - **Softmax** – the extension to multi-class; used usually for last layer
  - **But** *vanishing gradient* for very high or low input (i.e., slow or no further learning); non-zero centered output
  - **TanH**: variant for zero centered
- **ReLU**: computationally faster
  - **But** *dying ReLU* for negative or zero input (i.e., gradient is 0  $\rightarrow$  no learning)
  - Leaky ReLU: variant for negative or zero input to still have backpropagation

# Regularization

- Assumption: a model with smaller weights is less complex than one with larger weights
- Penalization for the weights added to the cost function
- $\lambda$  - regularization rate
- L1 norm is the sum of absolute values for the weights  $\lambda \sum |w_i|$
- L2 norm is the sum of squared values for the weights  $\lambda \sum w_i^2$

# Example of NN flow



# Package neuralnet in R

- Parameter *stepmax* denotes the maximum steps for training until stop.
- If the error *threshold* is not reached, then convergence is not achieved and no weights are given.
- If this happens, the default values of either *stepmax* = 1e+05 or *threshold* = 0.01 have to be increased.
- Function detects classification or regression based on the type of the variable to be predicted.

# NN in R

```
1 library(neuralnet)
2 library(datasets)
3 library(e1071)
4 library(caret)
5
6 data(iris)
7 dat <- iris
8
9 repeats <- 30
10 classColumn <- 5
11 accuracies <- vector(mode="numeric", length = repeats)
12 index <- 1:nrow(dat)
```

- Resilient backpropagation (Rprop) algorithm as the gradient descent algorithm for backpropagation
- Default activation function is logistic (if linear.output = FALSE)

# Scaling

- NN sensitive to scaling
- Scaling is not done by default in package neuralnet
- Repeat same principle, fit on the training set and then transform both training and test sets
  - Use function *preprocess* from library caret
  - Use method *range* for Min-Max scaler
- The function *predict* returns class probabilities for classification
  - Hence the maximum probability has to be taken to output the predicted class

# NN in R

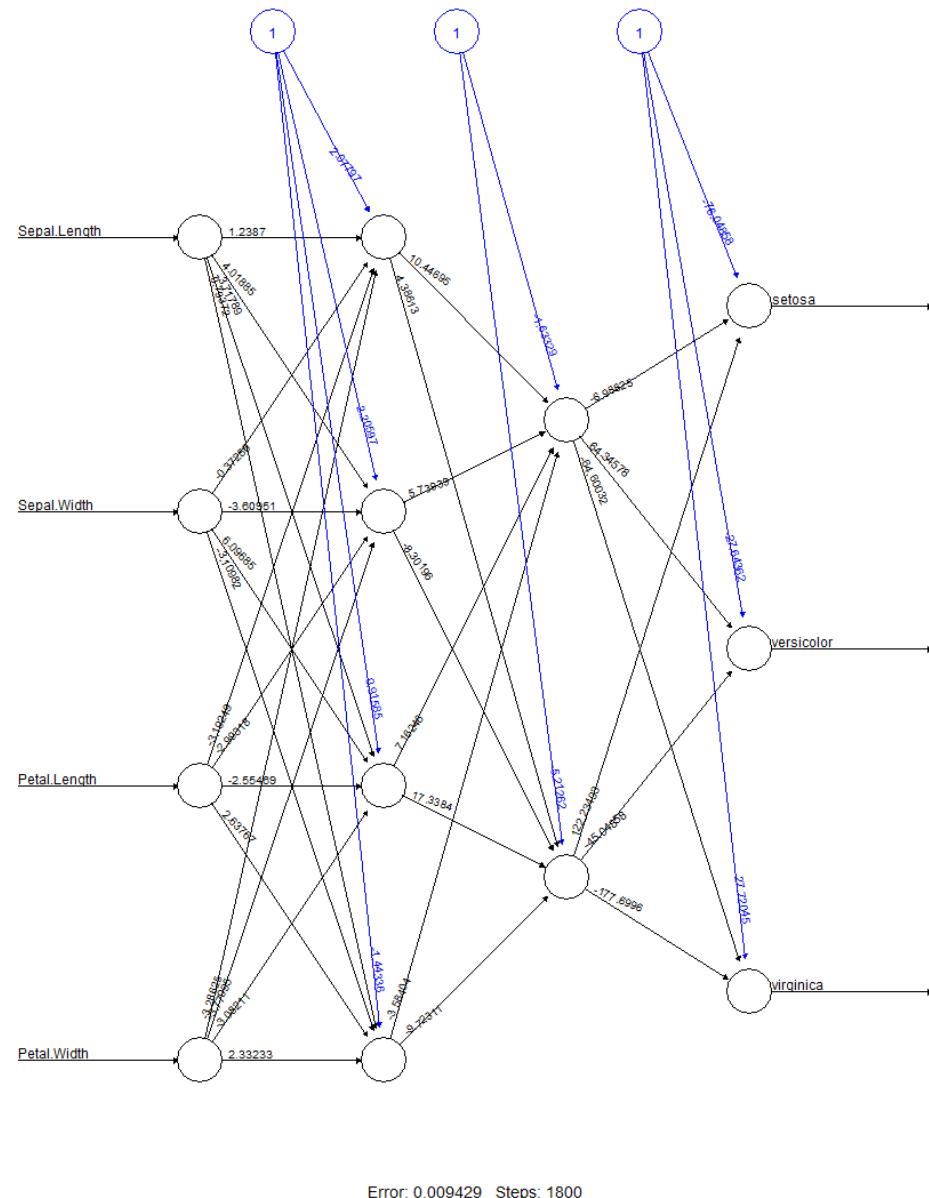
```
15 for (i in 1:repeats){  
16   testindex <- sample(index, trunc(length(index) / 3))  
17   testset <- dat[testindex, ]  
18   trainset <- dat[-testindex, ]  
19  
20   process <- preProcess(as.data.frame(trainset), method = c("range"))  
21   train_scaled <- predict(process, as.data.frame(trainset))  
22   test_scaled <- predict(process, as.data.frame(testset))  
23  
24   # Hidden specifies the number of neurons in each hidden layer  
25   # the length of the vector gives the number of hidden layers  
26  
27   # If activation function should not be applied to the output neurons,  
28   # linear_output = TRUE, otherwise FALSE.  
29   nn_model <- neuralnet(Species ~ ., data = train_scaled, hidden = c(4, 2), linear.output = FALSE, stepmax = 1e+06)  
30   nn_pred <- predict(nn_model, test_scaled[, -classColumn])  
31   labels <- c("setosa", "versicolor", "virginica")  
32   pred_label <- labels[max.col(nn_pred)]  
33   contab <- table(pred = pred_label, true = test_scaled[, classColumn])  
34   accuracies[i] <- classAgreement(contab)$diag  
35 }  
36  
37 # plot best representation of last model  
38 plot(nn_model)  
39  
40 print(accuracies)  
41 print(mean(accuracies))  
42 print(sqrt(var(accuracies)))  
43 print(summary(accuracies))  
44  
45 print("Confusion matrix of last run")  
46 print(contab)
```

# Classification results

```
[1] 0.96 0.96 0.98 0.94 0.98 0.98 0.98 0.96 0.94 0.92 0.98 0.96 0.96 1.00 0.94
[16] 0.98 0.98 0.94 0.96 0.92 0.96 0.96 0.96 0.98 0.88 0.96 0.90 0.94 0.94 0.96
[1] 0.9553333
[1] 0.02609444
      Min. 1st Qu. Median   Mean 3rd Qu.   Max.
0.8800 0.9400 0.9600 0.9553 0.9800 1.0000
[1] "Confusion matrix of last run"
      true
pred      setosa versicolor virginica
setosa     18       0       0
versicolor  0       13       0
virginica   0       2      17
```

# Plot NN

- Black arrows – variable weight values (without activation)
- Blue lines – bias values



# NN in Python

- In *sklearn* and the dedicated sublibrary *neural\_network*, there is the function *MLPClassifier*.
- Gradient algorithms:
  - Stochastic gradient descent (SGD)
  - Adam

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn import datasets
from sklearn.preprocessing import MinMaxScaler

repeats = 30
accuracies = []
scaler = MinMaxScaler()

ix, iy = datasets.load_iris(return_X_y=True)

for i in range(0, repeats):
    #random generation of training and test, 66%-33%
    ix_train, ix_test, iy_train, iy_test = train_test_split(ix, iy, test_size = 0.33)

    #scale data
    fit_scalar = scaler.fit(ix_train)
    ix_train_scaled = fit_scalar.transform(ix_train)

    nnm = MLPClassifier(hidden_layer_sizes=(4, 2), solver='adam', activation = 'tanh', max_iter = 10000)
    nnm.fit(ix_train_scaled,iy_train)

    ix_test_scaled = fit_scalar.transform(ix_test)
    iy_pred = nnm.predict(ix_test_scaled)
    accuracies.append(accuracy_score(iy_test, iy_pred))

print(accuracies)
```

# Results

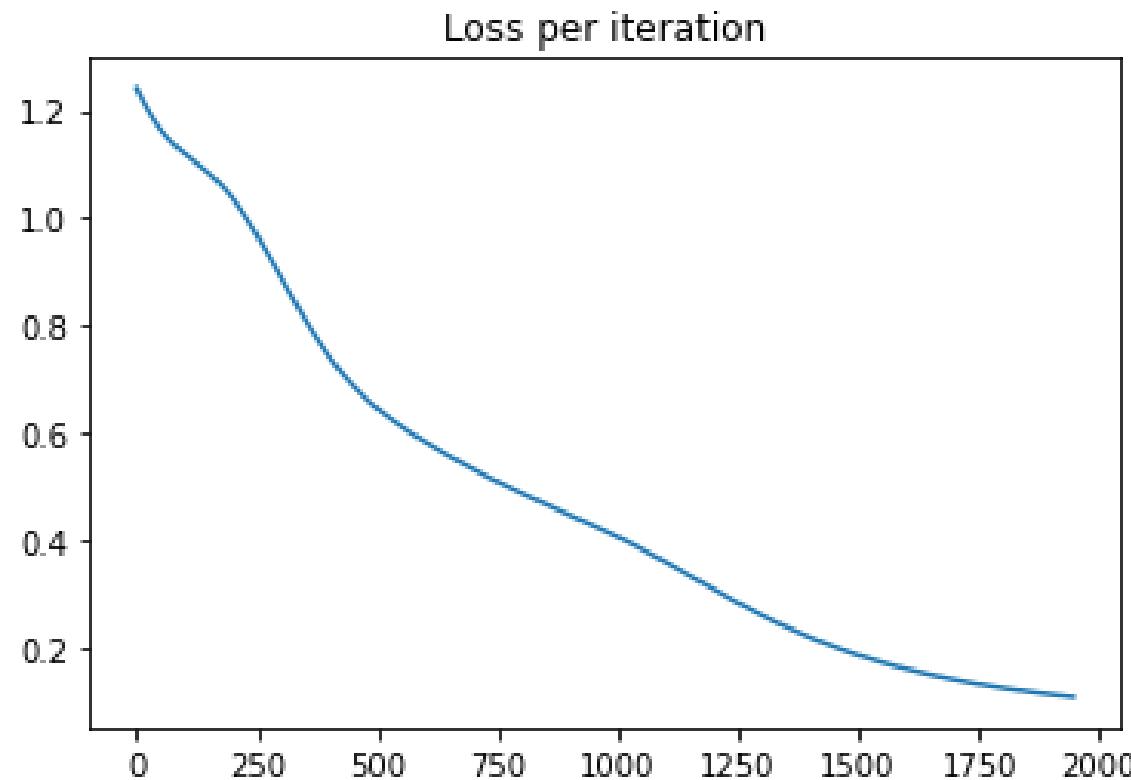
```
print("Mean accuracy", np.mean(accuracies))
print("Standard deviation", np.std(accuracies))
confusion_matrix(iy_test,iy_pred)
```

```
Mean accuracy 0.9606666666666668
Standard deviation 0.023371397523944144

array([[13,  0,  0],
       [ 0, 13,  3],
       [ 0,  0, 21]], dtype=int64)
```

# Loss decrease

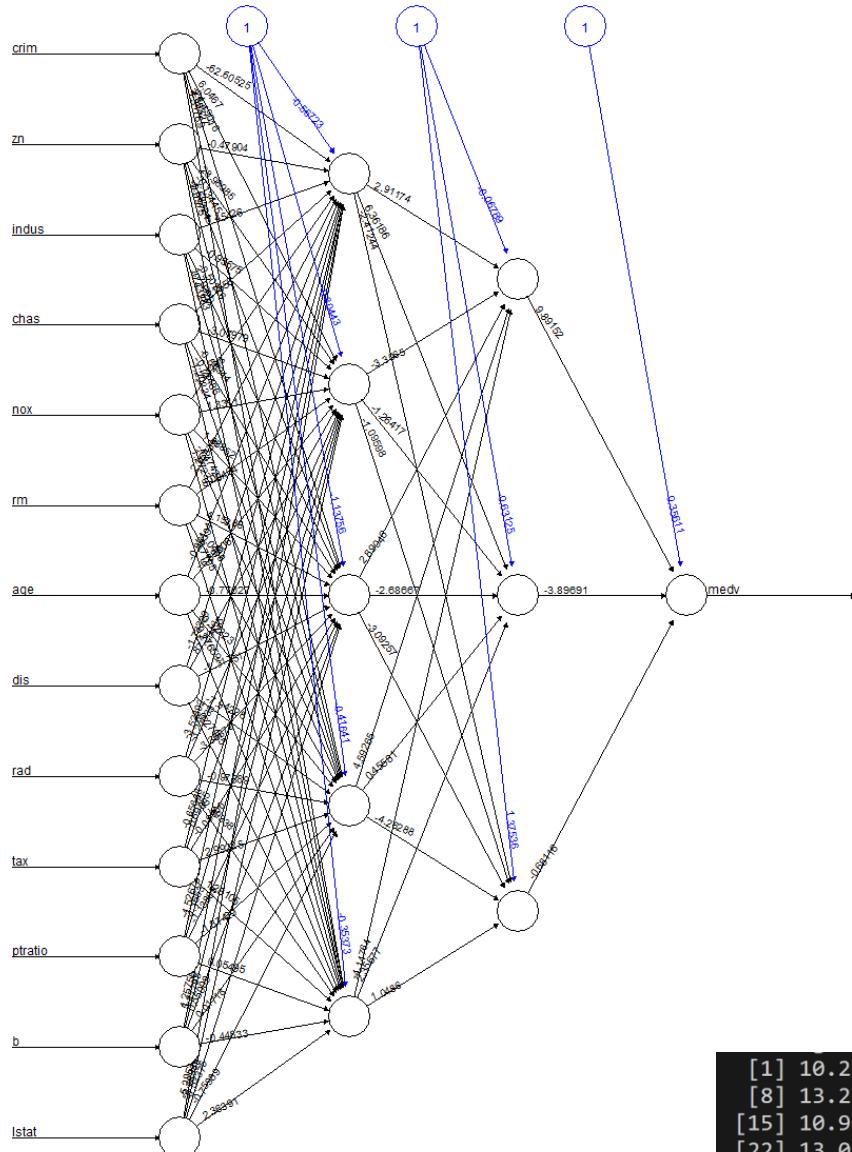
```
plt.plot(nn.loss_curve_)
plt.title("Loss per iteration")
```



# NN regression in R

```
1 library(neuralnet)
2 library(mlbench)
3 library(caret)
4
5 data(BostonHousing)
6
7 # transform factor columns to integer for neuralnet
8 dat <- data.frame(lapply(BostonHousing, as.numeric))
9
10 repeats <- 30
11 classColumn <- 14
12 mses <- vector(mode = "numeric", length = 30)
13 index <- 1:nrow(dat)
14
15 for (i in 1:repeats){
16   testindex <- sample(index, trunc(length(index) / 4))
17   testset <- dat[testindex, ]
18   trainset <- dat[-testindex, ]
19
20   dif <- max(trainset[classColumn]) - min(trainset[classColumn])
21
22   process <- preProcess(as.data.frame(trainset), method = c("range"))
23   train_scaled <- predict(process, as.data.frame(trainset))
24   test_scaled <- predict(process, as.data.frame(testset))
25
26   nn_model <- neuralnet(medv ~ ., data = train_scaled, hidden = c(5, 3), linear.output = FALSE, stepmax = 1e+06)
27   nn_pred <- predict(nn_model, test_scaled[, -classColumn])
28   mses[i] <- mean((dif * (nn_pred - test_scaled[, classColumn])))^2
29 }
30
31 plot(nn_model)
32
33 print(mses)
34 print(mean(mses))
35 print(sqrt(var(mses)))
```

# Results Boston housing



```
[1] 10.233588 10.128329 20.632011 31.248014 9.511252 9.487106 16.691137  
[8] 13.259778 21.159252 10.389862 22.701482 24.300544 12.448325 17.646360  
[15] 10.965124 10.858534 15.872975 21.218057 7.524402 22.229145 11.683342  
[22] 13.037122 22.236972 15.936926 7.215073 14.217332 8.488198 21.259411  
[29] 11.298770 9.324497  
[1] 15.10676  
[1] 6.058824
```

# NN regression in Python

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPRegressor
from sklearn import datasets
from sklearn.preprocessing import MinMaxScaler

repeats = 30
mses = []
r2s = []

bx, by = datasets.load_boston(return_X_y=True)

scaler = MinMaxScaler()

for i in range(0, repeats):
    #random generation of training and test, 75-25%
    bx_train, bx_test, by_train, by_test = train_test_split(bx, by, test_size = 0.25)

    fit_scalar = scaler.fit(bx_train)
    bx_train_scaled = fit_scalar.transform(bx_train)

    nnm = MLPRegressor(hidden_layer_sizes=(10, 10, 10), solver='adam', activation = 'relu', max_iter = 100000)
    nnm.fit(bx_train_scaled,by_train)

    bx_test_scaled = fit_scalar.transform(bx_test)
    by_pred = nnm.predict(bx_test_scaled)
    mses.append(mean_squared_error(by_test, by_pred))
    r2s.append(r2_score(by_test, by_pred))

print(mses)
print(r2s)
```

# Boston results

```
[17.492755007264936, 30.234505625912554, 10.02046071680119,  
19.037853523236794, 18.448766845564116, 20.87579946196899, 8.826236048369347,  
24.001715622499553, 18.878682321430265, 12.02416203522131,  
22.035370209214424, 9.329690404116555, 12.149204856579493,  
23.005473435020107, 8.454154531507836, 14.288911607219333,  
13.012196991202496, 10.710487549511683, 17.14418010662683,  
14.955289395359628, 23.199785268773645, 11.473790827049333,  
11.602084873366444, 11.009107901891278, 14.504382535062557,  
21.252839627988767, 10.808259854302339, 18.340272996999875,  
13.727357417027171, 14.585422802948548]
```

```
print("Mean MSE:", np.mean(mses))  
print("Standard deviation:", np.std(mses))  
print("R^2:", np.mean(r2s))
```

```
Mean MSE: 15.84764001333458  
Standard deviation: 5.318173954845035  
R^2: 0.8099998718540234
```

# Catastrophic forgetting

- When NN forget previously learned example after new data is learnt or when they are fine-tuned on more specific tasks.
  - Because of substantial weight re-adaptation to new examples
- Affecting large models more than smaller models
- Solutions:
  - Regularization: adding penalties for major weight adjustment
  - Progressive NN architecture (adding network parts for new tasks)
  - Ensembles
  - Rehearsal techniques (exposing the model to random old data)
  - Memory-augmented NN (NN + external memory storage)

# Curriculum learning

- At the data level:
  - Introduces training data from simple to complex examples
  - Question of deciding the order of the data
    - Given by expert (text, images)
    - According to model performance in time
    - Ensuring diversity
- At the model level
  - Augment the complexity of the architecture
    - Progressive model
    - Teacher-student
      - The student learns the task
      - The teacher determines the optimal learning parameters

# Grokking

- Delayed generalization
- Do models memorize or generalize?
- Models initially overfitting for a long time, but then achieving high performance on validation data
- Regularization (e.g. weight decay – L2 type) and learning rate seem to play a big role
- <https://pair.withgoogle.com/explorables/grokking/>

