

# Performance evaluation

Ruxandra Stoean

# Further bibliography

- N. Japkowicz, M. Shah, Evaluating Learning Algorithms, Cambridge University Press, 2011
- T. Hastie, R. Tibshirani, J. Friedman, The Elements of Statistical Learning, Springer Series in Statistics, 2001
- C. Huyen, Designing Machine Learning Systems: An Iterative Process for Production-Ready Applications, O'Reilly, 2022

# Directions

- Performance measures
  - Classification
  - Regression
- Error estimation
- Statistical significance

# Accuracy / Error rate

- The two complementary measures give an insight over the general performance.
- Accuracy is computed as the ratio between the amount of data labelled correctly and the total number of examples
  - The error rate is the ratio between the amount of samples labelled incorrectly and the total number of samples
- Disadvantages in the cases when:
  - Data with class imbalance
  - Different misclassification costs for each class

# Confusion matrix

- For  $k$  classes, a matrix of  $k \times k$ 
  - The element at position  $(i, j)$  gives the amount of data of class  $i$  that the classifier labels as  $j$
- Generically, for two classes (positive and negative labels)
- $N = TN + FP$ ,  $P = FN + TP$

	Predicted negative	Predicted positive
Real negative	True negative (TN)	False positive (FP)
Real positive	False negative (FN)	True positive (TP)

# Performance measures for a single class of interest 1/2

- Class of interest denoted by the positive.
- True-positive rate (TPR) gives the proportion of data of class  $i$  that are also labelled as  $i$  by the classifier.
- $$TPR = \frac{TP}{TP+FN}$$
- False-positive rate (FPR) – proportion of data that do not belong to class  $i$  but are labelled with that class.
- $$FPR = \frac{FP}{FP+TN}$$

# Performance measures for a single class of interest 2/2

- In the binary case, the two measures can be also defined for the opposite, negative class.
- True-negative rate  $TNR = \frac{TN}{TN+FP}$
- False-negative rate  $FNR = \frac{FN}{FN+TP}$
- TPR is also called **sensitivity** (or **recall**).
- TNR is also named **specificity**.

# Likelihood ratio

- Combines sensitivity and specificity to see the extent to which the classifier is efficient in prediction over the two classes.
- $LR_+ = \frac{Sensitivity}{1-Specificity} \rightarrow max$
- $LR_- = \frac{1-Sensitivity}{Specificity} \rightarrow min$
- $LR_+$  computes the times it is more likely that positive objects have a positive prediction compared to negative objects.
- $LR_-$  computes the times it is less likely that positive objects have a negative prediction compared to negative objects.

# LR in comparison

- Two algorithms A1 and A2.
- First,  $LR_+ \geq 1$ 
  - Contrarily, revert  $LR_+$  with  $LR_-$ .
- If  $LR_+(A1) > LR_+(A2)$  and  $LR_-(A1) < LR_-(A2)$ , then A1 is overall superior.
- If  $LR_+(A1) < LR_+(A2)$  and  $LR_-(A1) < LR_-(A2)$ , then A1 is superior in the confirmation of negative examples.
- If  $LR_+(A1) > LR_+(A2)$  and  $LR_-(A1) > LR_-(A2)$ , then A1 is superior in confirming positive examples.

# Positive predictive value (PPV)

- It is also called **precision**.
- The proportion of data that belong to class  $i$  of all detected by the algorithms as being class  $i$ .
- For the binary case:
- $$PPV = \frac{TP}{TP+FP}$$
- and the reverse NPV = 
$$\frac{TN}{TN+FN}$$
- The value  $PPV = a$  states that a positive prediction of the corresponding classifier will be true only in  $a\%$  of the cases.

# F1-score

- F-score combines precision and recall in a single measure, as the weighted harmonic mean of the two.
  - F1-score (balanced F-score): when the two measures have equal importance
- Useful for class imbalanced problems
- $Precision = \frac{TP}{TP+FP}$
- $Recall = \frac{TP}{TP+FN}$
- $F1 = \frac{2}{\frac{1}{Precision} + \frac{1}{Recall}} = \frac{2 \times Precision \times Recall}{Precision + Recall}$

# ROC analysis

- Receiver operating characteristic (**ROC**)
- **The ROC curve** for a classifier is a plot where
  - FPR on the horizontal axis
  - TPR on the vertical axis.
- ROC thus studies the relationship between sensitivity and specificity for a classifier.
- **The ROC space** – is a square of side 1.

# The ROC space 1/2

- The output of a classifier determines a point in the ROC space.
  - $(0,0)$  denotes a classifier that labels all data as negative ( $\text{TPR} = \text{FPR} = 0$ ).
  - $(1,1)$  denotes a classifier that labels all data as positive ( $\text{TPR} = \text{FPR} = 1$ ).
  - The classifiers whose output are positioned on the main diagonal ( $\text{TPR} = \text{FPR}$ ) label data in a random fashion.
  - Those whose output lies above the main diagonal have a better performance than random labelling.
    - Those whose output is situated under the diagonal have a performance weaker than random labelling.

# The ROC space 2/2

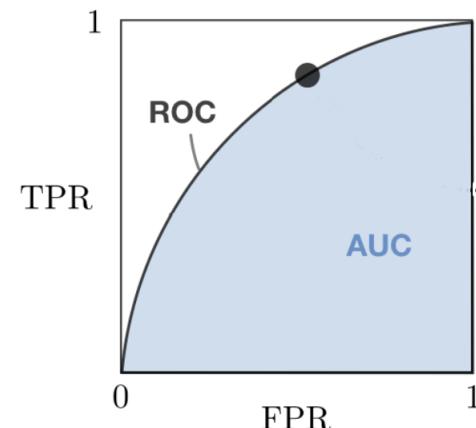
- For two points  $p_1$  and  $p_2$  in the ROC space,  $p_1$  is a better classifier than  $p_2$  if  $p_1$  is to the upper left side of  $p_2$ .
- $(1,0)$  ( $\text{FPR} = 1, \text{TPR} = 0$ ) denotes a classifier that has only wrong predictions.
- $(0,1)$  denotes the ideal classifier (labels all positive data correctly and does not make mistakes in labelling the negative examples).
- Classifiers with the output on the secondary diagonal have an equal performance both on the positive and the negative examples.

# The ROC curve

- An **operating point** corresponds to a decision threshold under which the classifier discriminates data as being positive or negative.
- Data with a score above this threshold are labelled as positive, and those with a score under are considered negative.
- By varying the threshold between the minimum and maximum score of the data, one such TPR and FPR are obtained that can be plotted:
  - This is the ROC curve.

# Area under the ROC curve (AUC)

- AUC – the measure of the performance of a classifier.
- The AUC value belongs to the interval  $[0, 1]$ 
  - 1 is the value for a perfect classifier.
  - The AUC value for a random classifier is 0.5.
  - For a reasonable performance, a classifier must have an AUC over 0.5.



# Statistical agreement measures

- Measures the weight of coincidence in the agreement between the classifier predictions and the true labels of the data.
- The most used statistics: **Cohen's k (kappa)**
  - Traditionally measures inter-rater reliability on qualitative ratings
  - Considers the possibility of random agreement
- $k = (p_0 - p_e) / (1 - p_e)$

# Calculation

	Pos	Neg
Pos	X	Y
Neg	Z	T

- $p_0$  – relative agreement observed between the two
  - The number of cases when both agree on positive and on negative labels over the total number of cases
  - $(X + T)/S$ ,  $S = X + Y + Z + T$
- $p_e$  – hypothetical probability of random agreement
  - (The proportion of cases when the classifier labeled as positive)  $X$  (the proportion of cases when the real data was positive) + (the proportion of cases when the classifier labeled negative)  $X$  (the proportion of cases when the data was negative)
  - $((X + Y)/S) * ((X + Z)/S) + ((Z + T)/S) * ((Y + T)/S)$

# Cohen's k

- $k = (p_0 - p_e) / (1 - p_e)$
- The k resulting value indicates:
  - $k < 0$  - "No agreement"
  - $0 < k < 0.2$  - "Slight agreement"
  - $0.2 < k < 0.4$  - "Fair agreement"
  - $0.4 < k < 0.6$  - "Moderate agreement"
  - $0.6 < k < 0.8$  - "Substantial agreement"
  - $0.8 < k < 1.0$  - "Almost perfect agreement"

# Metrics for regression

- $y_i$  - actual values,  $\hat{y}_i$  or  $f(\mathbf{x}_i)$  - predicted values,  $\bar{y}$  - mean of  $\mathbf{y}$
- $n$  - number of records,  $p$  – number of attributes
- Mean squared error (MSE)
  - Root mean squared error (RMSE)
- Mean absolute error (MAE) – absolute value of the difference
  - Useful for data with outliers
- Mean Absolute Percentage Error (MAPE)
  - Independent of the attribute scale

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

$$MAPE = \frac{1}{n} \sum_{i=1}^n \frac{|y_i - \hat{y}_i|}{y_i} \cdot 100\%$$

# Metrics for regression

- $y_i$  - actual values,  $\hat{y}_i$  or  $f(x_i)$  - predicted values,  $\bar{y}$  - mean of  $y$
- $n$  - number of records,  $p$  - number of attributes

$$R^2 = 1 - \frac{RSS}{TSS}$$

- Coefficient of determination  $R^2$

- RSS – residual sum of squares

$$RSS = \sum_{i=1}^n (y_i - f(x_i))^2$$

- TSS – total sum of squares

$$TSS = \sum_{i=1}^n (y_i - \bar{y})^2$$

- Adjusted  $R^2$  – takes into account the number of variables as well

$$R_{adj}^2 = 1 - \frac{(1 - R^2)(n - 1)}{n - p - 1}$$

# Error estimation 1/2

- It is used for an objective estimation of the performance of a learner.
- Data are split into three parts:
  - A training set from which the algorithm learns the associations between attribute values and outcomes.
  - A validation set to determine the error of prediction of the model (when the data collection is sufficiently large).
  - A test set to measure the generalization error.
  - For the last two sets, the data targets are considered to be unknown.



TRAINING

VALIDATION

TEST

# Error estimation 2/2

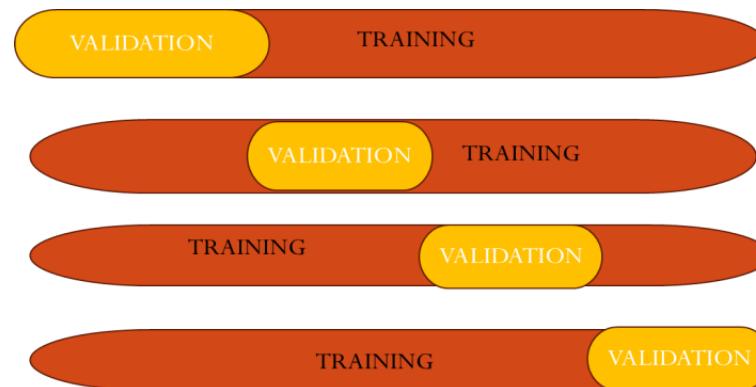
- **Cross-validation** estimates the prediction accuracy/error that the model will exhibit in practice.
- Training and test sets are selected a number of times following different schemes.
- The generalization ability of the learner is verified by computing the prediction accuracy/error on the test set as mean over several cross-validation runs.
- If the data set allows, the prediction accuracy/error is estimated on the validation set, before proceeding to the test set.

# Random subsampling

- The simplest **resampling** technique.
- $n$  runs of the learner are executed
  - Usually  $n = 30$ .
- In each run, the data set is split:
  - $2/3$  (or  $3/4$ ) the training set
  - $1/3$  (or  $1/4$ ) the test set
  - Data sampling into the two sets is done randomly without replacement.
- The final prediction accuracy/error is the mean of the corresponding values obtained in the 30 runs on the test set.

# K-fold cross-validation

- Split the training data set into  $k$  folds
  - Train on the  $k-1$  folds
  - Validate on the remaining fold
  - Average prediction accuracy/error on the  $k$  folds.



# Testing the statistical significance 1/2

- Useful (and compulsory) when:
  - The performance of a new model is compared to those of baseline algorithms.
  - The best algorithm for a given problem has to be selected.
- $n$  ( $n = 30$ ) runs of the algorithms through cross-validation are performed.

# Testing the statistical significance 2/2

- The null hypothesis  $H_0$  is presumed – the tested algorithms have a similar performance.
- The null hypothesis  $H_0$  is rejected if the p-value returned by the test is lower or equal to the significance level of 0.05.
- The safe option to test the statistical significance is the employment of a non-parametric test that does not make assumptions on the distribution of the performance values.

# Comparing learners

- Wilcoxon signed-rank test is a non-parametric method to compare 2 learners on 1 data set
  - An obtained p-value  $\leq 0.05$  implies that  $H_0$  is rejected
- Friedman test is also a non-parametric means to compare M machine learning algorithms over N datasets
  - Each learner is ranked per data set
  - The average rank is computed over all data sets
  - The Friedman statistics is calculated and interpreted
  - If the null hypothesis is rejected, a post-hoc Nemeyi test should be applied to find those that are statistically different from the others

# Model evaluation in R

```
1 library(e1071)
2 library(rpart)
3 library(mlbench)
4 library(fmsb) # for Cohen's k
5 library(ROCR) # for ROC
6 library(stats) # for Wilcoxon's test
7
8
9 data(PimaIndiansDiabetes)
10 dat <- PimaIndiansDiabetes
11
12 repeats <- 30
13 classColumn <- 9
14 accuraciesSVM <- vector(mode = "numeric", length = repeats)
15 accuraciesDT <- vector(mode = "numeric", length = repeats)
16 index <- 1:nrow(dat)
17
18 #cross-validation n = 30
19 for (i in 1:repeats){
20   testindex <- sample(index, trunc(length(index)/4))
21   testset <- dat[testindex, ]
22   trainset <- dat[-testindex, ]
23
24   # SVM model
25   svm.model <- svm(diabetes ~ ., data = trainset, kernel = "linear", cost = 1, probability = TRUE)
26   svm.pred <- predict(svm.model, testset[, -classColumn])
27
28   # DT model
29   rpart.model <- rpart(diabetes ~ ., data = trainset, method="class")
30   rpart.pred <- predict(rpart.model, testset[, -classColumn], type = c("class"))
```

```

32  # build confusion matrices
33  contabSVM <- table(pred = svm.pred, true = testset[, classColumn])
34  contabDT <- table(pred = rpart.pred, true = testset[, classColumn])
35
36  # collect accuracies
37  accuraciesSVM[i] <- classAgreement(contabSVM)$diag
38  accuraciesDT[i] <- classAgreement(contabDT)$diag
39 }
40
41 print("Accuracies SVM")
42 print(accuraciesSVM)
43 print("Accuracies DT")
44 print(accuraciesDT)
45
46 # mean accuracies
47
48 print("Mean accuracy SVM")
49 print(mean(accuraciesSVM))
50 print("Mean accuracy DT")
51 print(mean(accuraciesDT))
52
53 # standard deviation
54
55 print("StD SVM")
56 print(sqrt(var(accuraciesSVM)))
57 print("StD SVM")
58 print(sqrt(var(accuraciesDT)))
59
60 # confusion matrices
61 print("Confusion matrix SVM")
62 print(contabSVM)
63 print("Confusion matrix DT")
64 print(contabDT)
66 # Cohen's k predictions vs real output
67 print(Kappa.test(svm.pred, testset[, classColumn]))
68 print(Kappa.test(rpart.pred, testset[, classColumn]))

```

```
70 # ROC curve SVM
71
72 # set prediction through probabilities
73 # make probability = TRUE when training model
74 svm.probabilities <- predict(svm.model, testset[, -classColumn], probability = TRUE)
75
76 # extract column with probabilities for positive class
77 svm.pred <- attr(svm.probabilities, "probabilities")[, 1]
78
79 # take data real labels
80 # convert to integer and move to 1 for positive and 0 for negative
81 real <- testset[classColumn][,1]
82 labels <- as.integer(real) - 1
83
84 pred <- prediction(as.vector(svm.pred), labels)
85 perfSVM <- performance(pred, "tpr", "fpr")
86
87 print("AUC values SVM")
88 perf<- performance(pred, 'auc')
89 aucSVM <- as.numeric(perf@y.values)
90 print(aucSVM)
```

```
93 # ROC curve DT
94
95 # set prediction through probabilities
96 dt.probabilities <- predict(rpart.model, testset[, -classColumn], type = c("prob"))
97
98 # extract column with probabilities for positive class
99 dt.pred <- dt.probabilities[, 1]
100
101 # take data real labels
102 # convert to integer and move to 1 for positive and 0 for negative
103 real <- testset[classColumn][,1]
104 labels <- as.integer(real) - 1
105
106 pred <- prediction(as.vector(dt.pred), labels)
107 perfDT <- performance(pred, "tpr", "fpr")
108
109 print("AUC values DT")
110 perf <- performance(pred, 'auc')
111 aucDT <- as.numeric(perf@y.values)
112 print(aucDT)
```

# Results

```
[1] "Accuracies SVM"
[1] 0.7604167 0.7343750 0.7708333 0.8333333 0.7708333 0.7343750 0.7552083
[8] 0.7812500 0.7656250 0.7812500 0.8177083 0.7604167 0.7864583 0.7812500
[15] 0.7864583 0.7916667 0.7968750 0.7604167 0.7968750 0.7500000 0.7760417
[22] 0.7864583 0.7552083 0.7812500 0.8177083 0.7395833 0.8177083 0.7343750
[29] 0.8020833 0.8020833
[1] "Accuracies DT"
[1] 0.7343750 0.7552083 0.7604167 0.7812500 0.7760417 0.7395833 0.7239583
[8] 0.7291667 0.7552083 0.7291667 0.7552083 0.7447917 0.7656250 0.7291667
[15] 0.7760417 0.7291667 0.7239583 0.7604167 0.7552083 0.6822917 0.6875000
[22] 0.7187500 0.7916667 0.7187500 0.7760417 0.7187500 0.7500000 0.7500000
[29] 0.7656250 0.7395833
[1] "Mean accuracy SVM"
[1] 0.7776042
[1] "Mean accuracy DT"
[1] 0.7440972
[1] "StD SVM"
[1] 0.02645682
[1] "StD SVM"
[1] 0.0259421
[1] "Confusion matrix SVM"
      true
pred  neg  pos
  neg 112   23
  pos   15   42
[1] "Confusion matrix DT"
      true
pred  neg  pos
  neg 103   26
  pos   24   39
```

# Cohen's K agreement

```
Estimate Cohen's kappa statistics and test the null hypothesis that  
the extent of agreement is same as random (kappa=0)
```

```
data: svm.pred and testset[, classColumn]  
Z = 6.611, p-value = 1.909e-11  
95 percent confidence interval:  
 0.4146655 0.6741318  
sample estimates:  
[1] 0.5443987
```

```
$Judgement  
[1] "Moderate agreement"
```

```
$Result
```

```
Estimate Cohen's kappa statistics and test the null hypothesis that  
the extent of agreement is same as random (kappa=0)
```

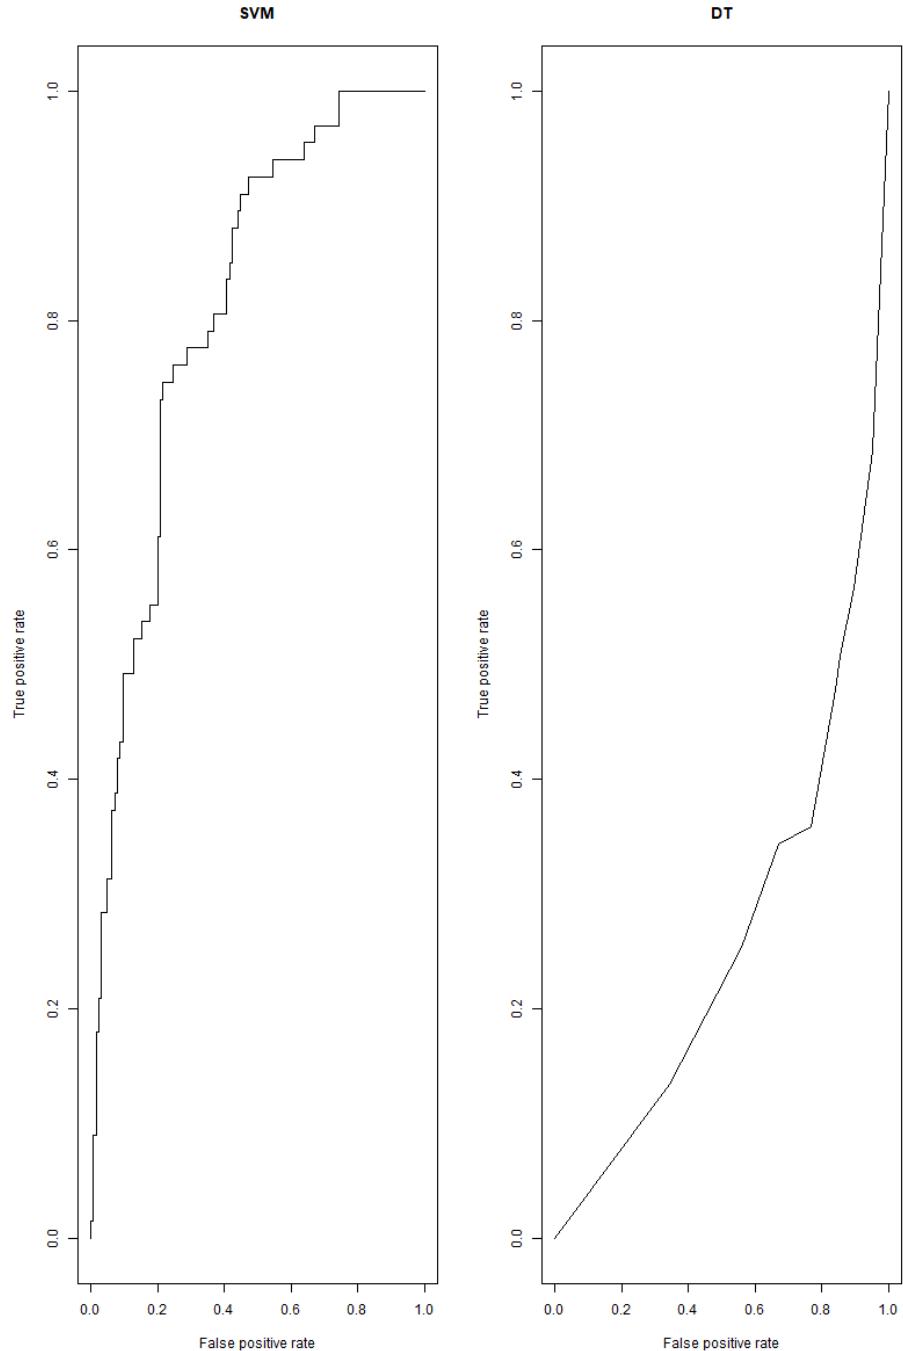
```
data: rpart.pred and testset[, classColumn]  
Z = 5.1332, p-value = 1.425e-07  
95 percent confidence interval:  
 0.2744795 0.5537885  
sample estimates:  
[1] 0.414134
```

```
$Judgement  
[1] "Moderate agreement"
```

# ROC curve

```
par(mfrow = c(1, 2))
plot(perfSVM)
title('SVM')
plot(perfDT)
title('DT')
```

```
[1] "AUC values SVM"
[1] 0.8103881
[1] "AUC values DT"
[1] 0.2671642
```



# Wilcoxon test results

```
print(wilcox.test(accuraciesSVM, accuraciesDT))
```

```
Wilcoxon rank sum test with continuity correction

data: accuraciesSVM and accuraciesDT
W = 738.5, p-value = 1.991e-05
alternative hypothesis: true location shift is not equal to 0
```

- There is statistical difference between the performances of the two classifiers.

# Python

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import cohen_kappa_score
from sklearn.metrics import f1_score
from sklearn import metrics

repeats = 30

accuraciesSVM = []
f1SVM = []

accuraciesDT = []
f1DT = []

data = pd.read_csv("diabetes.csv")

#take first n-1 columns for x and last column for y
dx = data.iloc[:, :-1]
dy = data.iloc[:, -1]

for i in range(0, repeats):
    #random generation of training and test, 75-25%
    dx_train, dx_test, dy_train, dy_test = train_test_split(dx, dy, test_size = 0.25)
```

```
# SVM training
svmm = svm.SVC(kernel='linear', probability = True)
svmm.fit(dx_train,dy_train)
dy_pred = svmm.predict(dx_test)
accuraciesSVM.append(accuracy_score(dy_test, dy_pred))
f1SVM.append(f1_score(dy_test, dy_pred))

# DT training
dtm = DecisionTreeClassifier()
dtm.fit(dx_train,dy_train)
dy_pred2 = dtm.predict(dx_test)
accuraciesDT.append(accuracy_score(dy_test, dy_pred2))
f1DT.append(f1_score(dy_test, dy_pred2))

print("Accuracies SVM")
print(accuraciesSVM)
print("Accuracies DT")
print(accuraciesDT)

print("Mean accuracy SVM", np.mean(accuraciesSVM))
print("Standard deviation SVM", np.std(accuraciesSVM))
print("Mean F1-score SVM", np.mean(f1SVM))

print("Confusion matrix SVM")
confusion_matrix(dy_test,dy_pred)

print("Cohen's k SVM")
print(cohen_kappa_score(dy_test,dy_pred))

print("Mean accuracy DT", np.mean(accuraciesDT))
print("Standard deviation DT", np.std(accuraciesDT))
print("Mean F1-score DT", np.mean(f1DT))

print("Confusion matrix DT")
confusion_matrix(dy_test,dy_pred2)

print("Cohen's k DT")
print(cohen_kappa_score(dy_test,dy_pred2))
```

# Results

```
Accuracies SVM  
[0.796875, 0.7760416666666666, 0.8072916666666666, 0.78125, 0.7604166666666666, 0.75, 0.796875, 0.7692708333333333, 0.71875]  
Accuracies DT  
[0.661458333333334, 0.708333333333334, 0.6666666666666666, 0.6875, 0.692708333333334, 0.71875, 0.765625, 0.765625, 0.765625]  
Mean accuracy SVM 0.7692708333333333  
Standard deviation SVM 0.024911228387386245  
Mean F1-score SVM 0.6304878352031682  
Confusion matrix SVM  
Cohen's k SVM  
0.45495836487509467  
Mean accuracy DT 0.7036458333333333  
Standard deviation DT 0.0312022204182295  
Mean F1-score DT 0.5757284697002113  
Confusion matrix DT  
Cohen's k DT  
0.2269399707174231
```

```
print("Confusion matrix DT")  
confusion_matrix(dy_test,dy_pred2)  
✓ 0.0s
```

Confusion matrix DT

```
array([[86, 37],  
       [29, 40]], dtype=int64)
```

```
print("Confusion matrix SVM")  
confusion_matrix(dy_test,dy_pred)  
✓ 0.0s
```

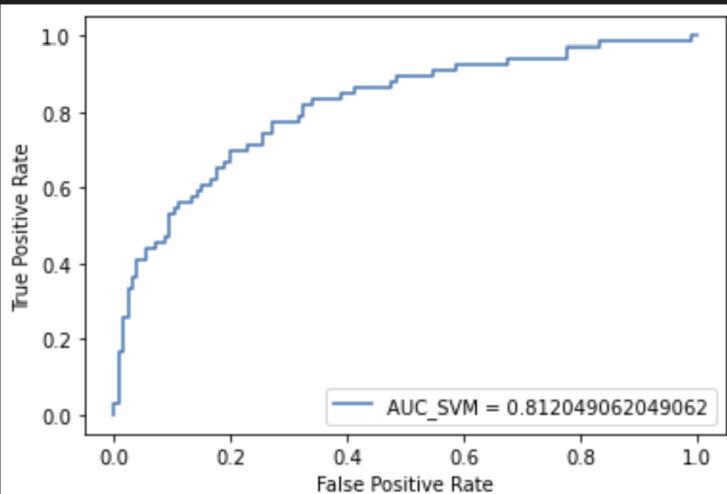
Confusion matrix SVM

```
array([[108, 15],  
       [28, 41]], dtype=int64)
```

```
# plot the ROC curve for last run of SVM
dy_pred_prob = svmm.predict_proba(dx_test)[:,1]
fpr, tpr, _ = metrics.roc_curve(dy_test, dy_pred_prob)
auc = metrics.roc_auc_score(dy_test, dy_pred_prob)

plt.plot(fpr,tpr,label = "AUC_SVM = " + str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```

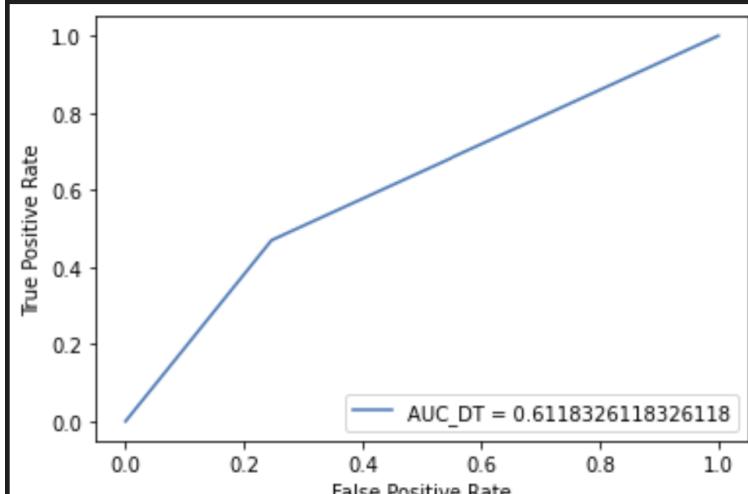
✓ 0.1s



```
# plot the ROC curve for last run of DT
dy_pred_prob2 = dtm.predict_proba(dx_test)[:,1]
fpr, tpr, _ = metrics.roc_curve(dy_test, dy_pred_prob2)
auc = metrics.roc_auc_score(dy_test, dy_pred_prob2)

plt.plot(fpr,tpr,label = "AUC_DT = " + str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```

✓ 0.1s



```
import scipy.stats as stats
```

```
stats.wilcoxon(accuraciesSVM, accuraciesDT)
```

✓ 0.0s

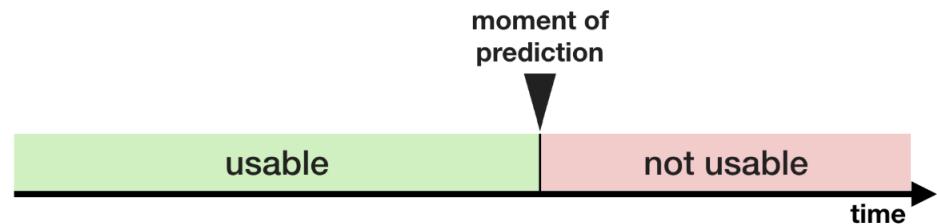
```
WilcoxonResult(statistic=0.0, pvalue=1.714847401033251e-06)
```

# Data leakage

- When training data contains information that will not be found in deployment.
- Train-test contamination
  - When you perform a data pre-processing action before splitting into training-test-validation
  - e.g. data imputation of missing values, scaling

# Target leakage

- Predictors that were added to the data set after the moment of prediction



[kaggle.com/learn](https://kaggle.com/learn)

Age	Fever	Headache	Target	
			Bacterial infection	Antibiotic
34	Y	Y	Y	Y
30	Y	N	N	N
23	N	Y	N	N



# Homework 1/2

- Take one problem (Wisconsin, car or customer) and models from last homework (NN, DT, Ensembles)
- If the task is classification, for each model:
  - Conduct cross-validation by random subsampling with 30 runs.
  - Obtain mean accuracy, standard deviation, F1-score and the confusion matrix.
  - Compute the agreement between predictions and real targets with Cohen's k.
  - Obtain the ROC curve and the AUC value.
  - Apply the Wilcoxon test to compare the predictions of the most accurate two classifiers.

# Homework 2/2

- If the task is regression, for each model:
  - Conduct cross-validation cross-validation by random subsampling with 30 runs.
  - Obtain MSE, RMSE, MAE, MAPE,  $R^2$  and Adjusted  $R^2$ .
  - Apply the Wilcoxon test to compare the predictions of the least erroneous two regressors.
  - Cohen's k cannot be applied, as the data is not qualitative.
- **(Optional)** Construct at least 3 classifiers for the Iris, Wisconsin breast cancer and Wine data sets and try the application of Friedman's test to rank their performance.



Using  
accuracy to  
evaluate the model

---

Using  
ROC-AUC to  
evaluate the model