

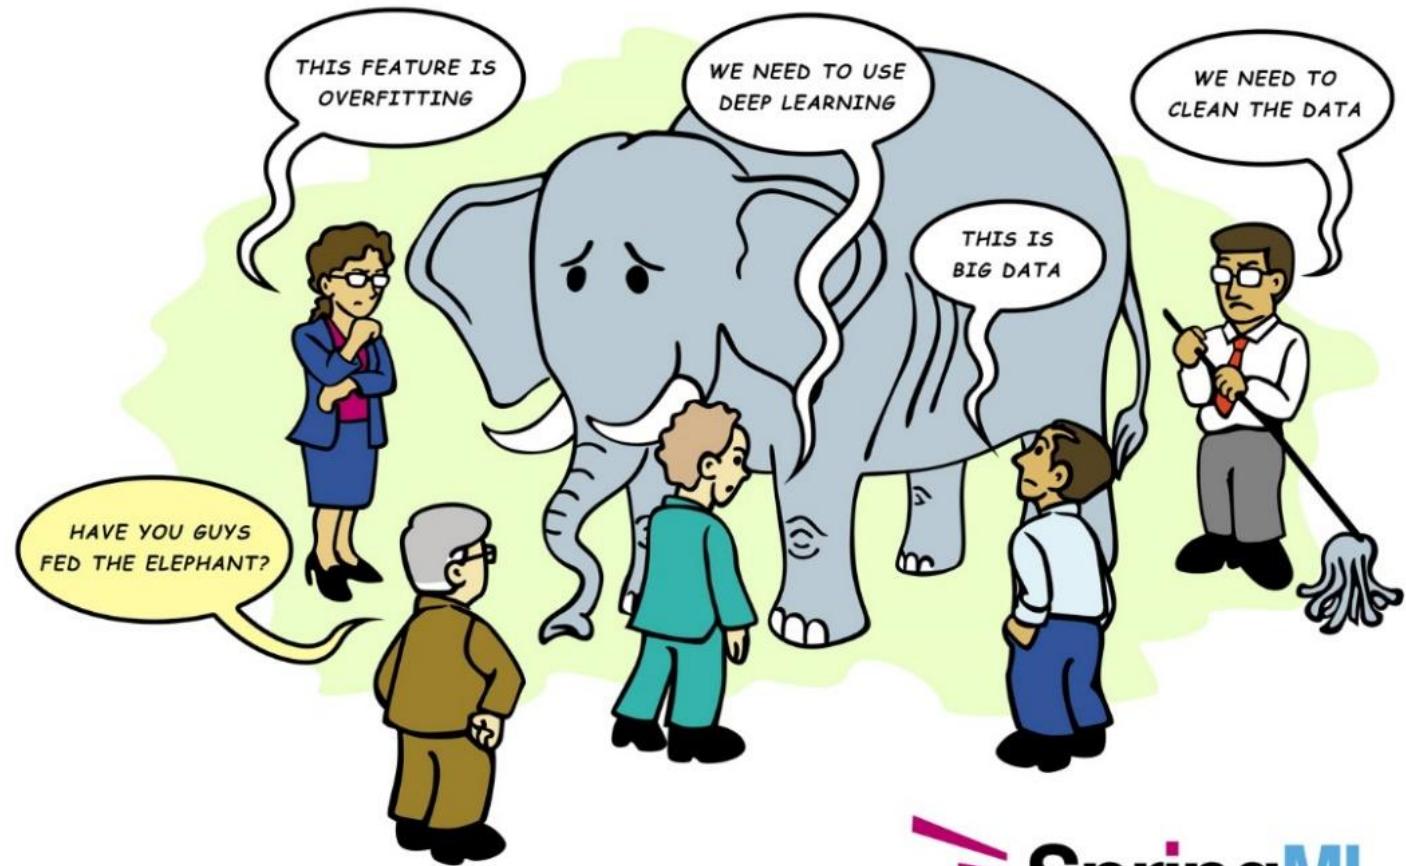
Deep learning for image processing



Ruxandra Stoean

Further bibliography

- Ian Goodfellow et al, Deep Learning (Adaptive Computation and Machine Learning series), MIT Press, 2016 <http://www.deeplearningbook.org/>
- John Kelleher, Deep Learning (MIT Press Essential Knowledge series), 2019
- Charu Aggarwal, Neural Networks and Deep Learning: A Textbook 2nd ed, Springer, 2023
- Aston Zhang et al, Dive into Deep Learning, Cambridge University Press, 2023
- Simon Prince, Understanding Deep Learning, MIT Press, 2023
- Jonah Gamba, Deep Learning Models, A Practical Approach for Hands-On Professionals, 2024

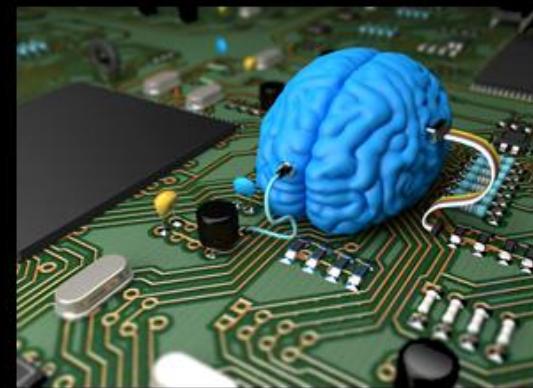


 SpringML

Deep Learning



What society thinks I do



What my friends think I do



What other computer
scientists think I do



What mathematicians think I do



What I think I do

```
In [1]:  
import keras  
Using TensorFlow backend.
```

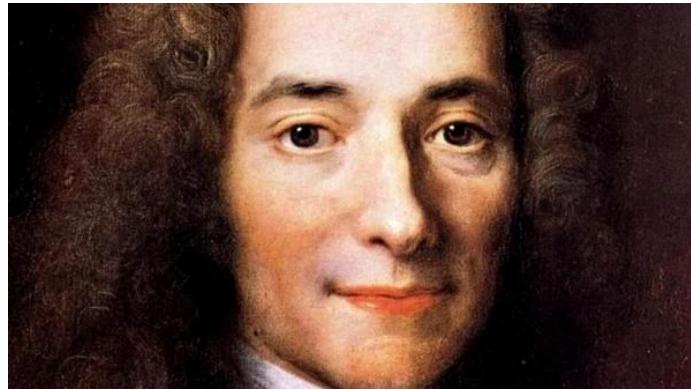
What I actually do

Definitions

- “For most flavors of the old generations of learning algorithms … performance will plateau. … deep learning … is the first class of algorithms … that is scalable. … performance just keeps getting better as you feed them more data.” - Andrew Ng
- “The hierarchy of concepts allows the computer to learn complicated concepts by building them out of simpler ones. If we draw a graph showing how these concepts are built on top of each other, the graph is deep, with many layers. For this reason, we call this approach to AI deep learning.” – Ian Goodfellow
- “Deep learning [is] … a pipeline of modules all of which are trainable. … deep because [has] multiple stages in the process of recognizing an object and all of those stages are part of the training” - Yann LeCun
- “At which problem depth does Shallow Learning end, and Deep Learning begin? Discussions with DL experts have not yet yielded a conclusive response to this question. [...], let me just define for the purposes of this overview: problems of depth > 10 require Very Deep Learning.”- Jurgen Schmidhuber



"I never guess. It is a capital mistake to theorize before one has data. Insensibly one begins to twist facts to suit theories, instead of theories to suit facts."



“The more I read, the more I acquire,
the more certain I am that I know nothing.”

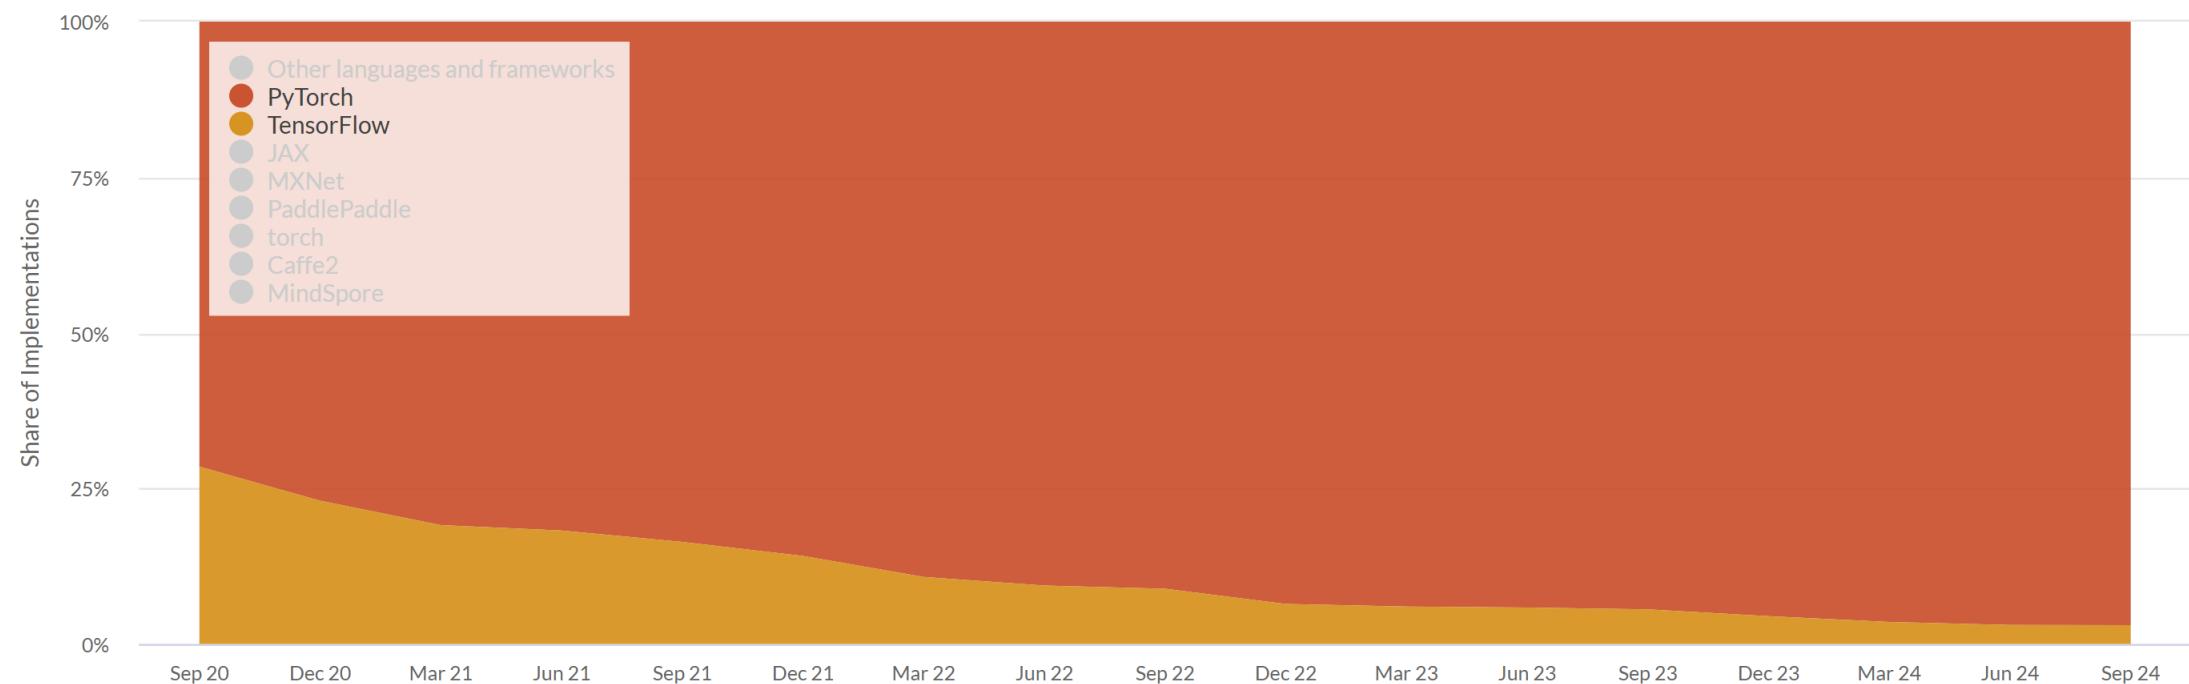
— Voltaire

Frameworks for implementing deep learning

Tensorflow	Pytorch
Developed by Google	Developed by Meta
Visualization with Tensorboard	Efficient memory usage
Simple high-level API	Python low-level coding
Mature library	Younger library, but increasing in popularity

Tensorflow vs. Pytorch

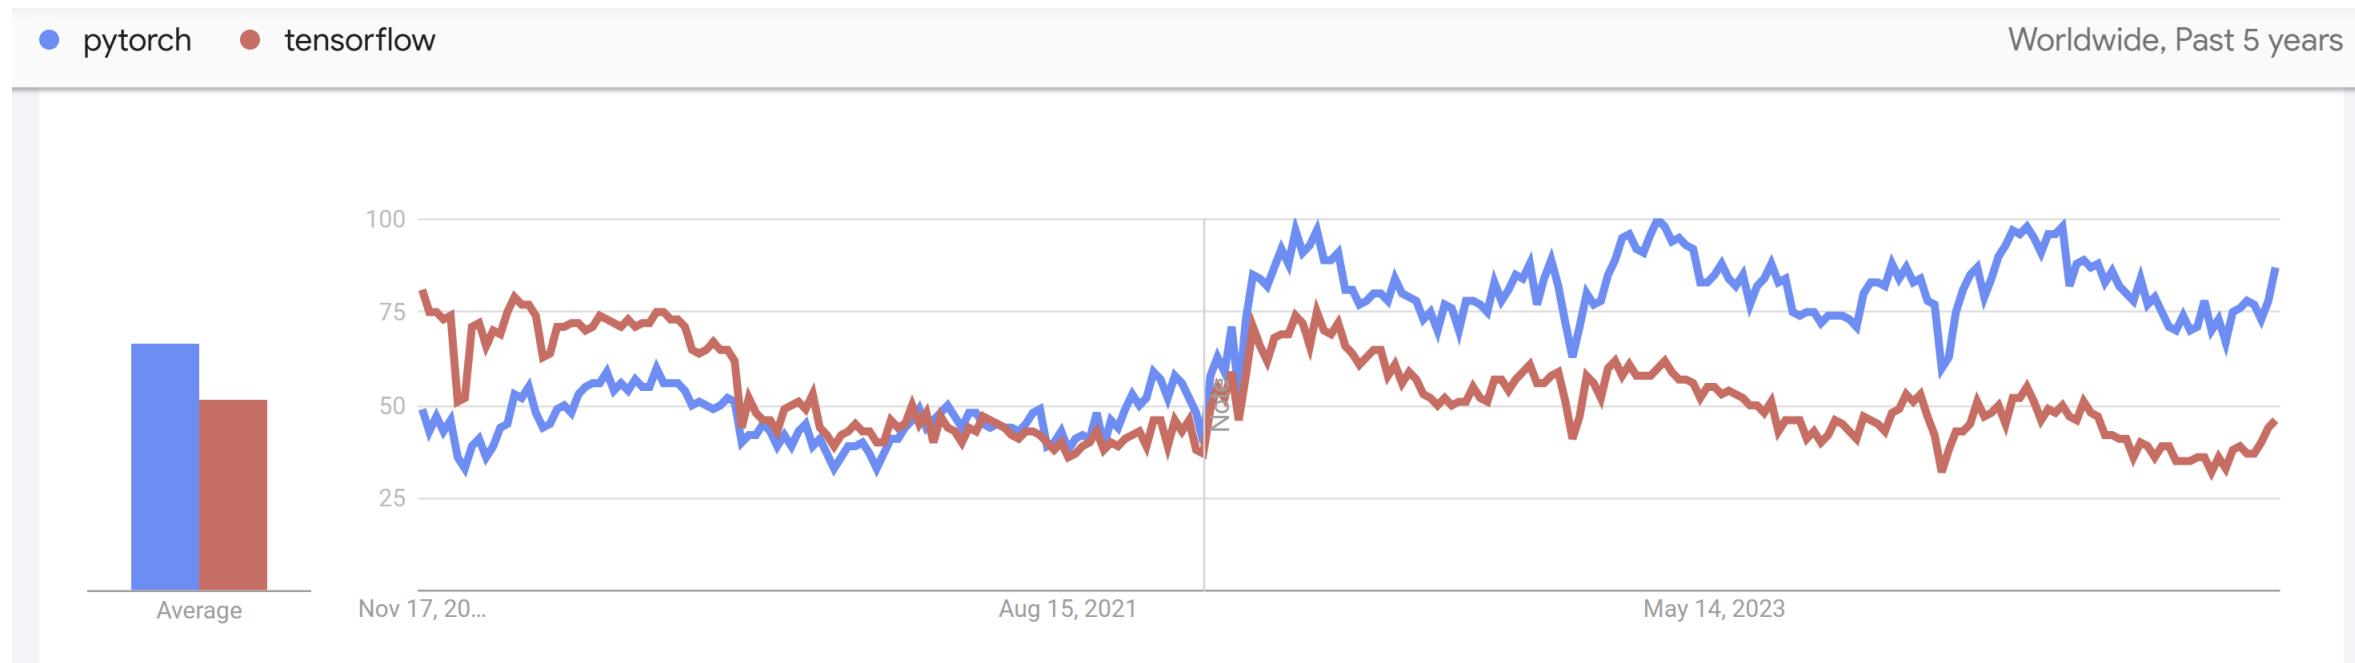
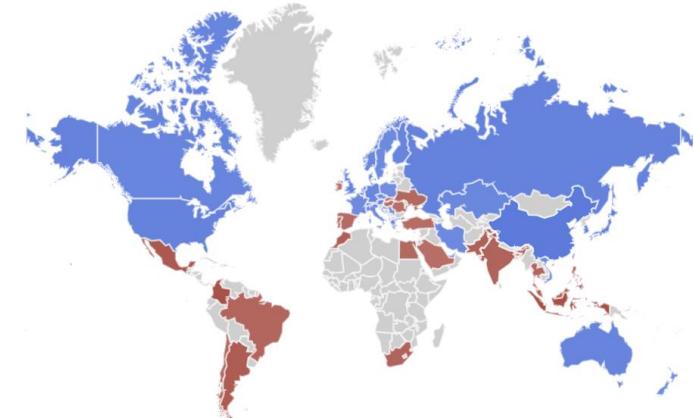
Paper Implementations grouped by framework



<https://paperswithcode.com/trends>

● pytorch ● tensorflow

Tensorflow vs. Pytorch



<https://trends.google.com/trends/explore?date=today%205-y&q=pytorch,tensorflow&hl=en>

Working with Tensorflow/Keras under R

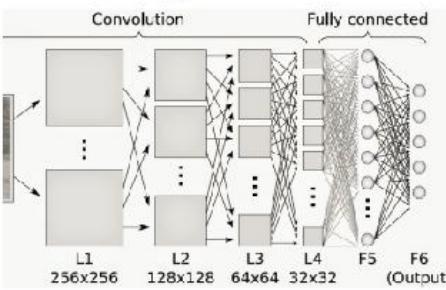
- Installation:
 - Install Rtools last version from <https://cran.r-project.org/bin/windows/Rtools/>
 - `Install.packages("dplyr")`
 - `Library(devtools)`
 - `Install_github("rstudio/reticulate")`
 - `Install.packages("keras")` for Tensorflow and Keras

Image processing. Convolutional neural networks

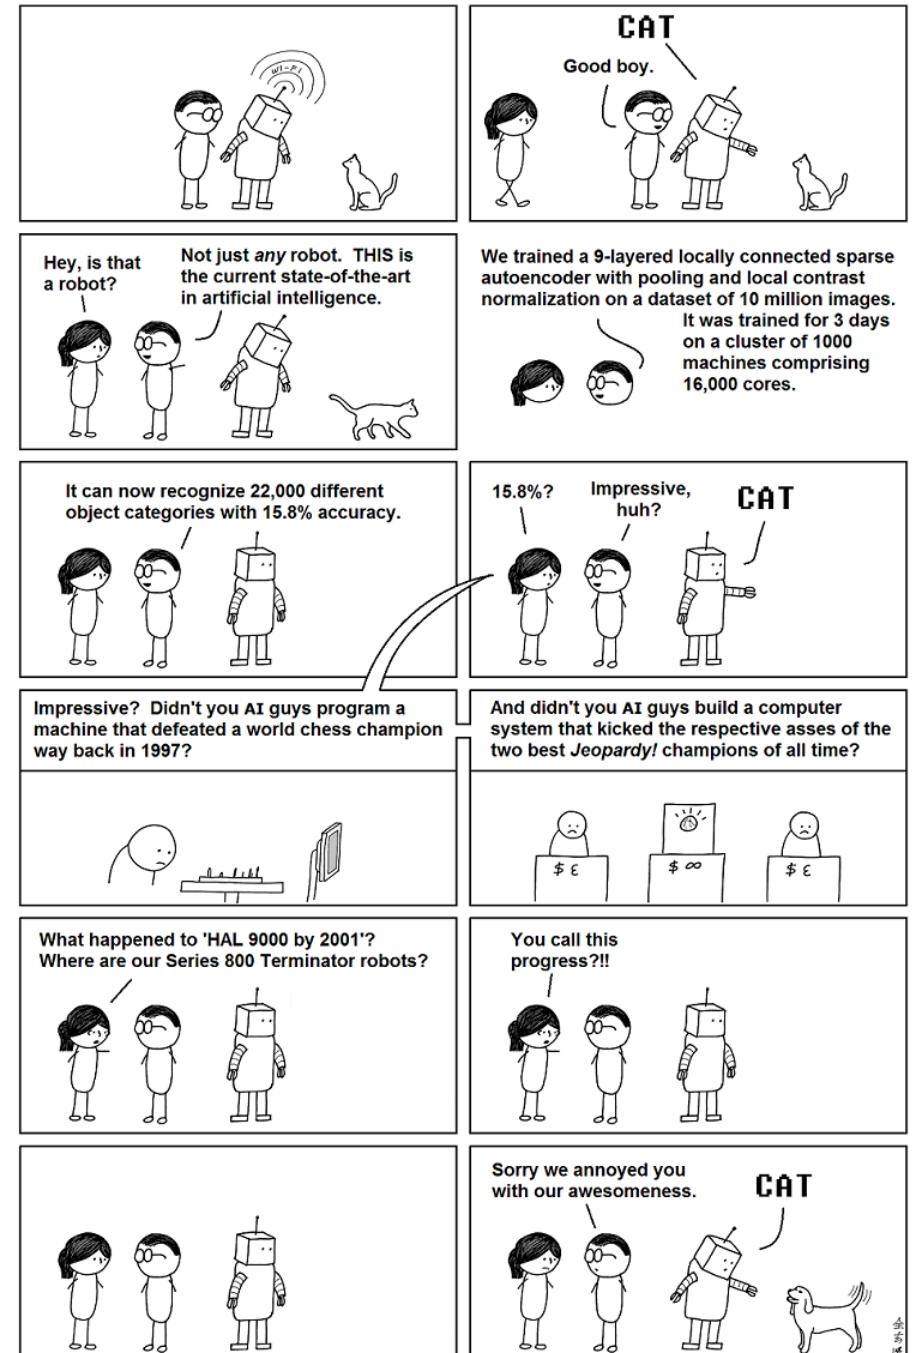
Convolutional neural networks (CNN)



what he's actually referring to

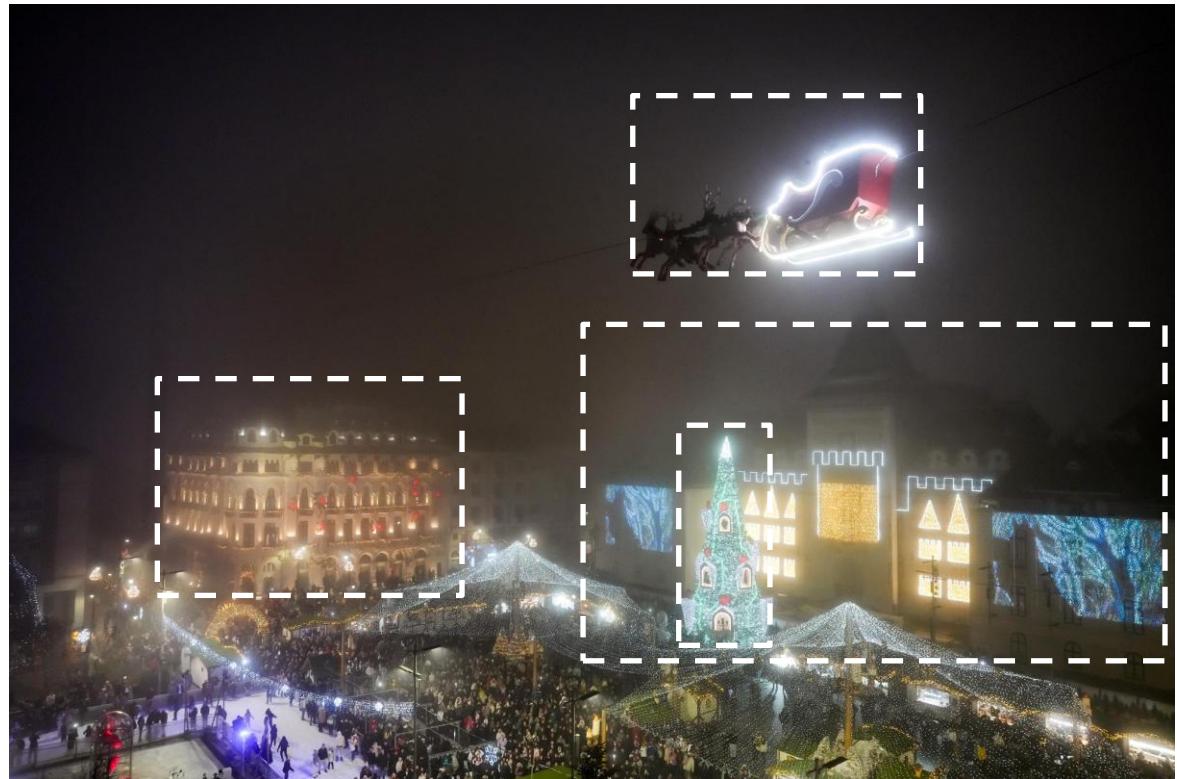


<https://www.quora.com/What-are-some-good-data-science-statistics-machine-learning-jokes>



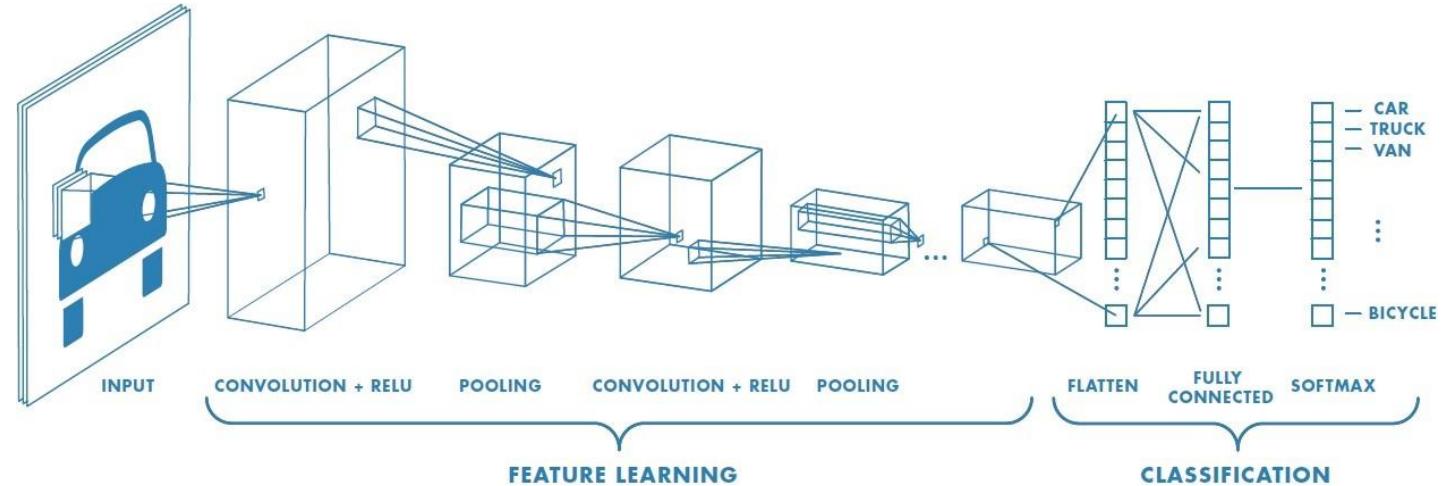
CNN

- Automatic feature learning: from low-level to high-level
- Important applications in computer vision
 - **Classification** – There is a building in this image
 - **Semantic Segmentation** – These are the pixels of the building
 - **Object Detection** – There are buildings in this image
 - **Instance Segmentation** – There are several instances of buildings in this image and these are the pixels of each



Architecture

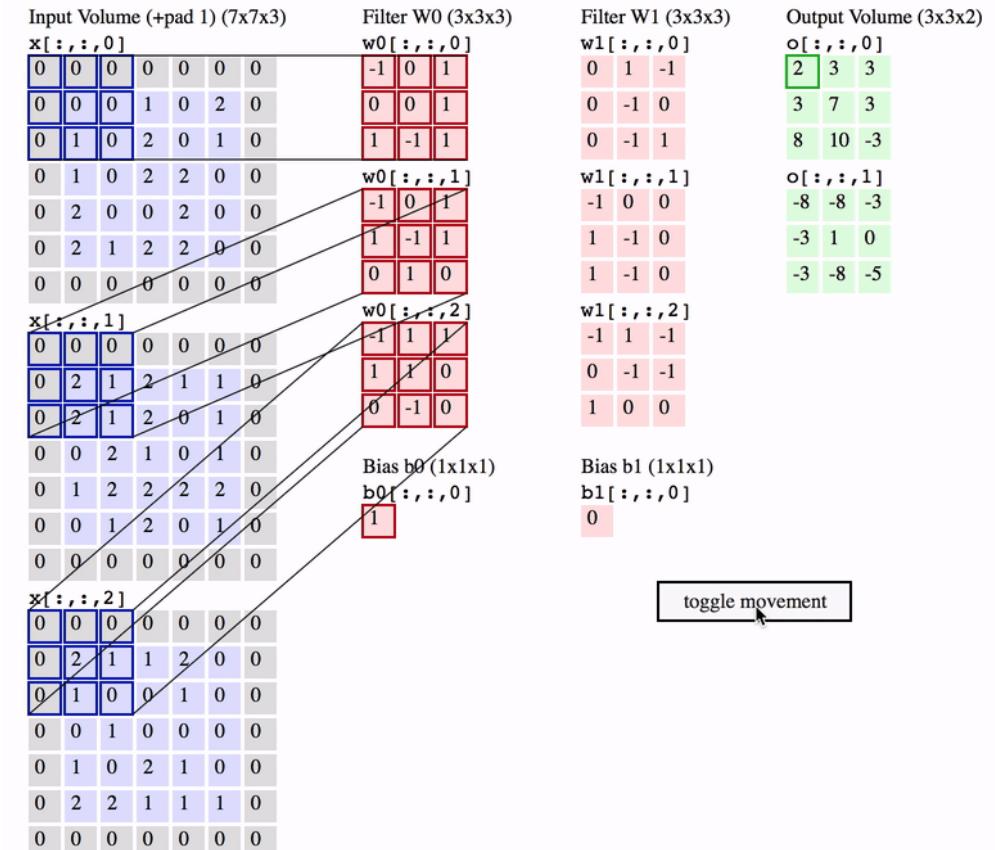
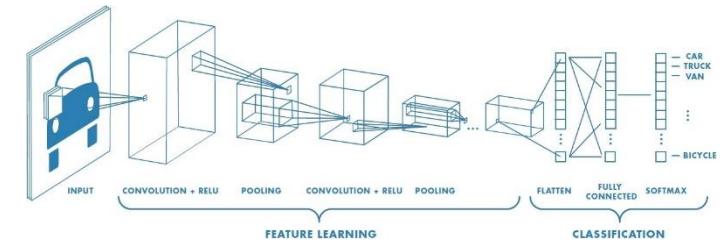
- Layers
 - Feature learning
 - Convolution
 - ReLU
 - Pooling
 - Classification
 - Fully connected



<https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f148>

Convolution

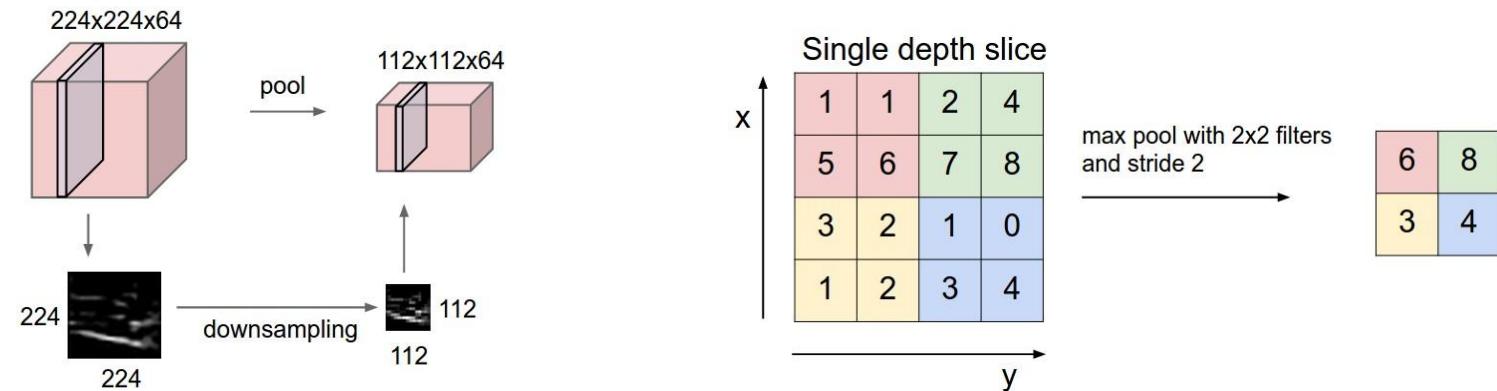
- Input
 - Volume of Width x Height x Depth
- Kernel (or filter) - a shared set of weights
- Forward pass – convolution between filter and input volume
- Result: KD activation maps of size KSxKS stacked in a new volume
- Four hyperparameters:
 - Kernel size (KS)
 - Kernel depth (number of filters) (KD)
 - Stride
 - Zero-padding



<http://cs231n.github.io/convolutional-networks/>
<https://www.youtube.com/watch?v=AQirPKrAyDg>

ReLU & Pooling

- Rectified Linear Unit – transfer layer for nonlinearity
- (Max) Pooling – downsamples the volume by taking the (maximum) value
 - Window size
 - Stride



<http://cs231n.github.io/convolutional-networks/>
<https://www.youtube.com/watch?v=AQirPKrAyDg>

The practice

- Choice of the proper architecture
- Parameter dependence on the problem
- Large runtime
- Great computing power required
- Small sample size in real-world
- Overfitting
- Difficult model interpretation
- Not plug & play!

"Just one more episode"



"Just one more page"



"Just one more piece"



"Just one more layer"



Machine Learning Memes for Convolutional Teens

DEEP LEARNING

EXPECTATIONS:

`import tensorflow`



PROBLEM SOLVED!

REALITY:

DAY 1 Iteration 1
Epoch 1/5
[.....]



DAY 3 Iteration 2
Epoch 2/5
[==>.....]



DAY 8 Iteration 6
Epoch 4/5
[=====]>.....]



DAY 15 Iteration 10
Epoch 5/5
[=====]>.....]



Memory Error: out of memory



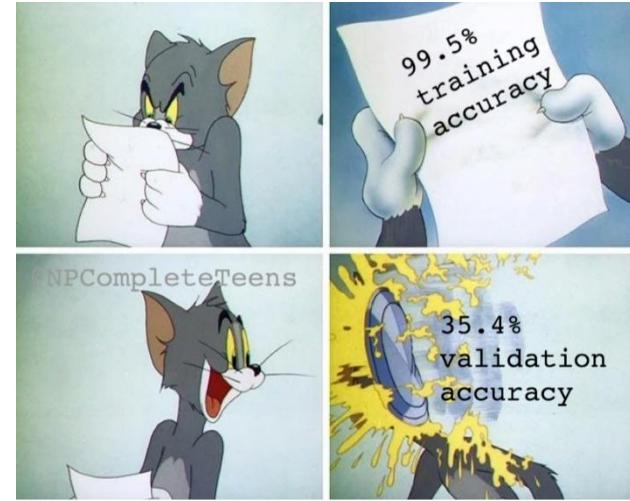
<https://medium.com/@elmira/deep-learning-frustration-4c26d69609f2>

Parametrization

- Convolutional kernels
 - Sizes, depths, strides
 - Pooling layers
 - Window sizes, strides
 - Dropout rates
 - Batch size
 - Number of epochs
 - Initial weights
 - Optimizers
 - Learning rate
 - Activation functions
 - Number of units in the fully connected layers
 - Topology
-
- Tuning
 - Manual
 - Automatic

Overfitting

- When the model is too complex for the data
- Working with real-world problems
 - Small sample size
- Ways to combat:
 - Data augmentation
 - Flip, rotate, scale, crop, translate, Gaussian noise
 - Generative adversarial networks (GAN): one network generates, the other evaluates
 - Dropout layer
 - Regularization
 - Weight penalty (decay) L1 and L2
 - Early stopping
 - Use of checkpoints to save the model each epoch
 - Pick the best candidate from validation result after last epoch



Machine Learning Memes for Convolutional Teens

Image classification

MNIST digits classification in Keras/TF under R

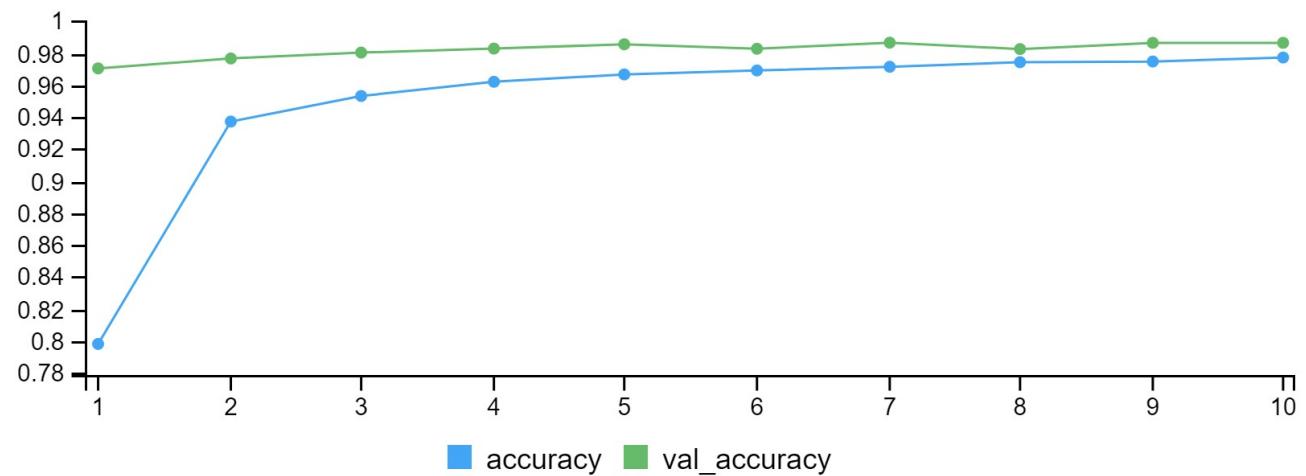
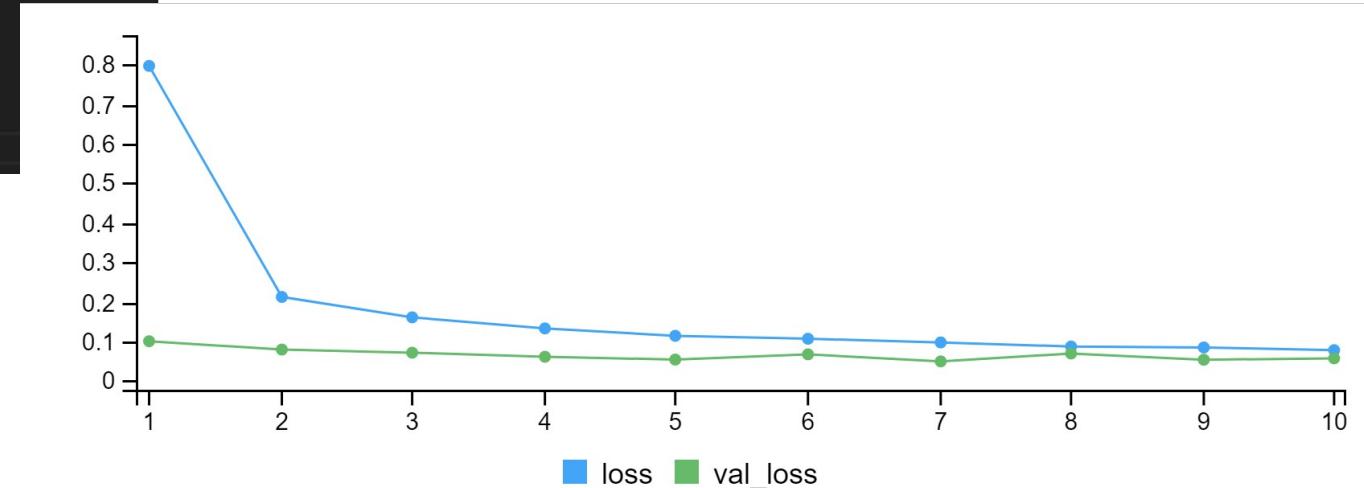
- A 2-layer CNN for MNIST digit recognition

```
1 library("keras")
2
3 # load MNIST data set
4 mnist <- dataset_mnist()
5 c(x_train, y_train) %<-% mnist$train
6 c(x_test, y_test)  %<-% mnist$test
7 rbind(dim(x_train), dim(x_test))
8
9 # transform to format for CNN (length, width, number of color channels)
10 x_train_original <- x_train
11 x_test_original <- x_test
12 x_train           <- array_reshape(x_train, c(nrow(x_train), 28, 28, 1))
13 x_test            <- array_reshape(x_test, c(nrow(x_test), 28, 28, 1))
14 input_shape       <- c(28, 28, 1)
15
16 # transform classes codified by 0-9 digits to one-hot binary encoding
17 y_train_original <- y_train
18 y_test_original <- y_test
19 y_train          <- to_categorical(y_train, 10)
20 y_test           <- to_categorical(y_test, 10)
```

```
23 # construct CNN architecture
24 # 1. Feature Learning
25 # two convolutionary layers, one with 8 filters of 3 x 3, the second with 16 filter of size 5 x 5
26 # RELU nonlinearity
27 # Max Pooling of size 2 x 2 and stride 2
28 # Dropout of 0.25
29 # 2. Classification
30 # one fully connected layer with 100 units and ReLU activation
31 # Dropout of 0.5
32 # Last fully connected layer to the 10 target classes with softmax activation
33
34 model <- keras_model_sequential() %>%
35   layer_conv_2d(filters = 8, kernel_size = c(3,3), activation = 'relu', padding="same",
36                 input_shape = input_shape) %>%
37   layer_conv_2d(filters = 16, kernel_size = c(5,5), activation = 'relu', padding="same") %>%
38   layer_max_pooling_2d(pool_size = c(2, 2)) %>%
39   layer_dropout(rate = 0.25) %>%
40   layer_flatten() %>%
41   layer_dense(units = 100, activation = 'relu') %>%
42   layer_dropout(rate = 0.5) %>%
43   layer_dense(units = 10, activation = 'softmax')
44
45 summary(model)
46
47 # compile model with categorical cross-entropy loss, Adam optimizer and accuracy as metric
48 model %>% compile(
49   loss = 'categorical_crossentropy', # for multi-class classification
50   optimizer = 'adam', # optimizer
51   metrics = c('accuracy') # accuracy as model performance
52 )
```

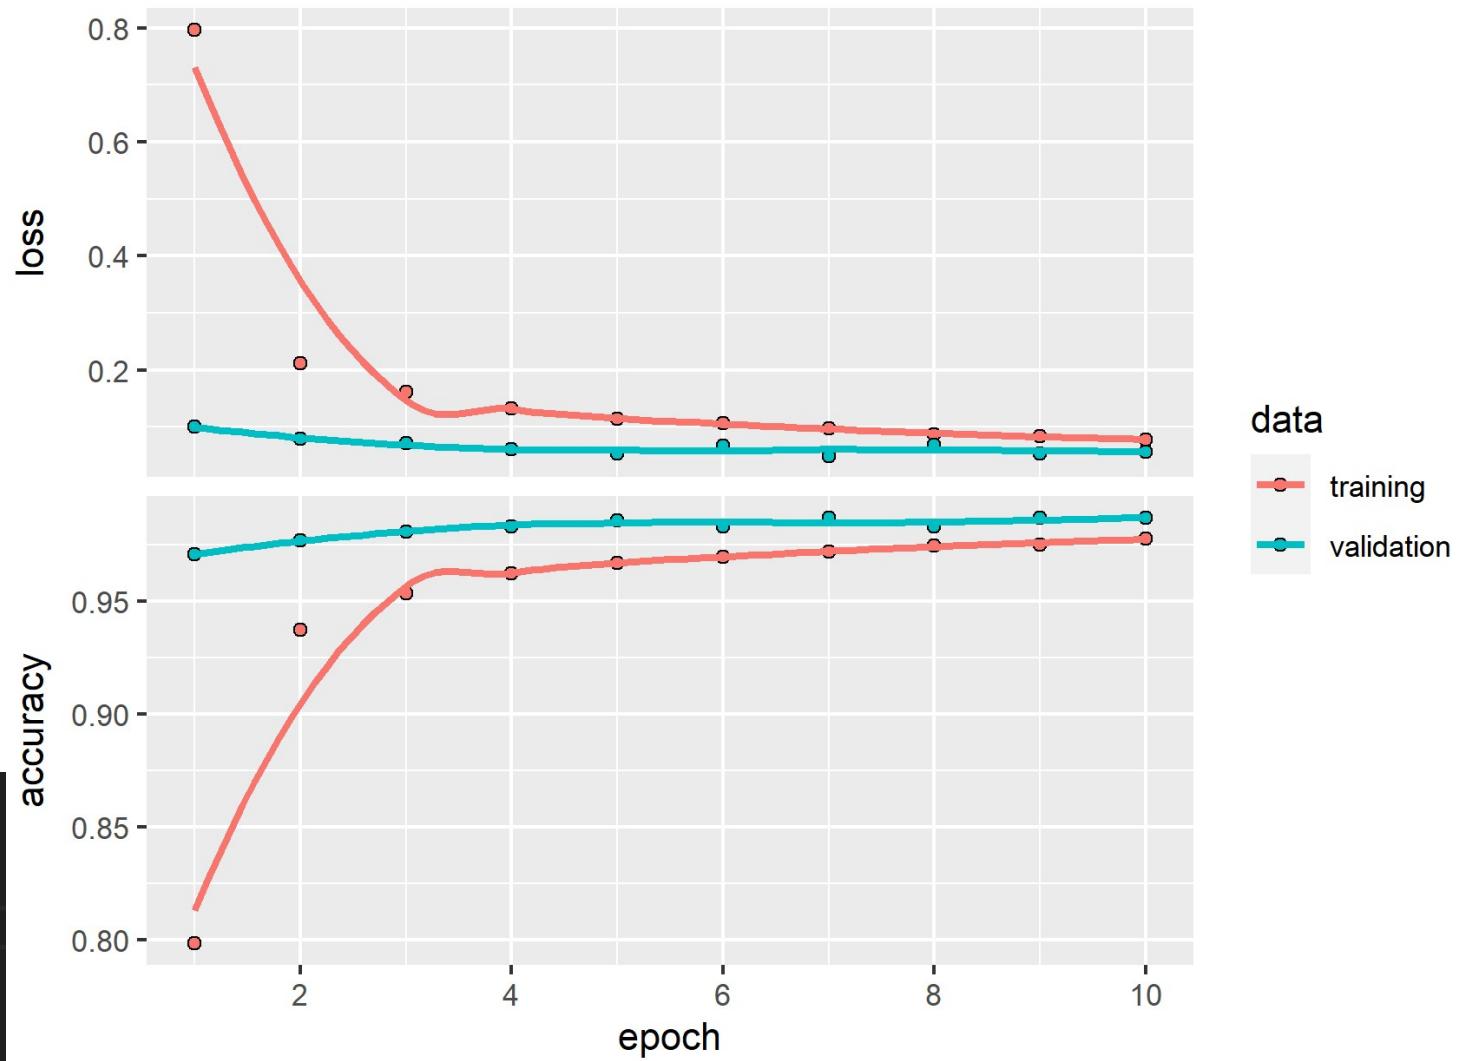
Real-time training history

```
54 # train model
55 history <- model %>% fit(
56   x_train, y_train,
57   epochs = 10,           # number of epochs (10 complete passes of the training set)
58   batch_size = 25,       # size of batch in one data pass
59   validation_split = 0.25 # training 75% - validation 25%
60 )
61
62 # plot loss and accuracy along training
63 plot(history)
```



Plot training history

```
62 # plot loss and accuracy along training
63 plot(history)
64
65 # evaluate model on test
66 model %>% evaluate(x_test, y_test)
67
68 # CNN model predictions as class probabilities
69 # label with class of highest probability
70
71 y_test_predicted_val <- model %>% predict(x_test)
72 y_test_predicted_ten <- y_test_predicted_val %>% k_argmax()
73 y_test_predicted <- as.numeric(y_test_predicted_ten)
```



Visualize test predictions

```
77 par(mfcol=c(6,6))
78 par(mar=c(0, 0, 3, 0), xaxs='i', yaxs='i')
79 for (idx in 1:36) {
80   im <- x_test_original[idx,]
81   im <- t(apply(im, 2, rev))
82   if (y_test_predicted[idx] == y_test_original[idx]) {
83     color <- '#008800'
84   } else {
85     color <- '#bb0000'
86   }
87   image(1:28, 1:28, im, col=gray((0:255)/255),
88         xaxt='n', main=paste0(y_test_predicted[idx], " (", y_test_original[idx], ")"),
89         col.main=color)
90 }
```

Accuracy:
98.79%

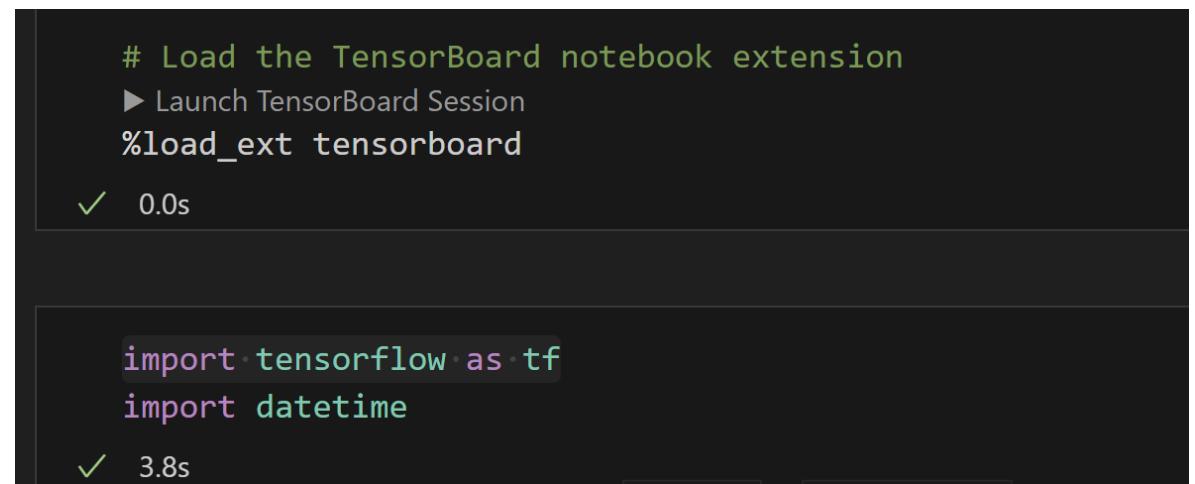
7 (7)	4 (4)	9 (9)	5 (3)	4 (4)	3 (3)
7	4	9	5	4	3
2 (2)	9 (9)	0 (0)	4 (4)	0 (0)	1 (1)
2	8	0	4	0	1
1 (1)	5 (5)	1 (1)	9 (9)	7 (7)	3 (3)
1	5	1	9	7	3
0 (0)	9 (9)	5 (5)	6 (6)	4 (4)	4 (4)
0	9	5	6	4	4
4 (4)	0 (0)	9 (9)	6 (6)	0 (0)	7 (7)
4	0	9	6	0	7
1 (1)	6 (6)	7 (7)	5 (5)	1 (1)	2 (2)
1	6	7	5	1	2

CNN for CIFAR-10 in Tensorflow/Keras

- <https://www.cs.toronto.edu/~kriz/cifar.html>
- A data set of 60000 images, size 32x32
- Real-world pictures
- 10 classes, 6000 images per class
- 50000 training data, 10000 test data

TensorBoard

- Real time and storables measurements/visualizations during the machine learning flow
- !Delete contents of “\AppData\Local\Temp\.tensorboard-info“ to be able to launch new TB session!
- See TB directly in notebook



The image shows a Jupyter Notebook interface with two code cells. The top cell contains the following code:

```
# Load the TensorBoard notebook extension
▶ Launch TensorBoard Session
%load_ext tensorboard
```

The output of this cell shows a green checkmark and the time "0.0s". The bottom cell contains the following code:

```
import tensorflow as tf
import datetime
```

The output of this cell shows a green checkmark and the time "3.8s".

```
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
import numpy as np

# load CIFAR10 data set

(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# normalize pixel values
train_images, test_images = train_images / 255.0, test_images / 255.0
```

```
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

# Plot some samples from the training data

plt.rcParams.update({'text.color': "black",
                     'axes.labelcolor': "white"})

plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()
```

```
# The CNN model

# The convolutional layers for feature learning
# input_shape for pictures of the form [length, width, number of channels]
# CNN architecture: 3 convolutional layers, depths 32, 64, 64 and sizes 3x3, max-pooling with size 2x2
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.Dropout(0.5))

# The flatten layer collapses the spatial dimensions of the input into 1D
model.add(layers.Flatten())

# The dense layers for classification: one fully connected layer with 64 neurons and the final dense layer with 10 outputs for the classes
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation = 'softmax'))

# Summary of the CNN architecture
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 6, 6, 64)		0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
dropout (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 64)	65600
dense_1 (Dense)	(None, 10)	650
<hr/>		
Total params: 122,570		
Trainable params: 122,570		
Non-trainable params: 0		

```
▶ Launch TensorBoard Session  
%tensorboard --logdir logs
```

[6] ✓ 3.6s

```
# view training/validation history in real time  
from tensorboard import notebook  
notebook.list()  
✓ 0.0s
```

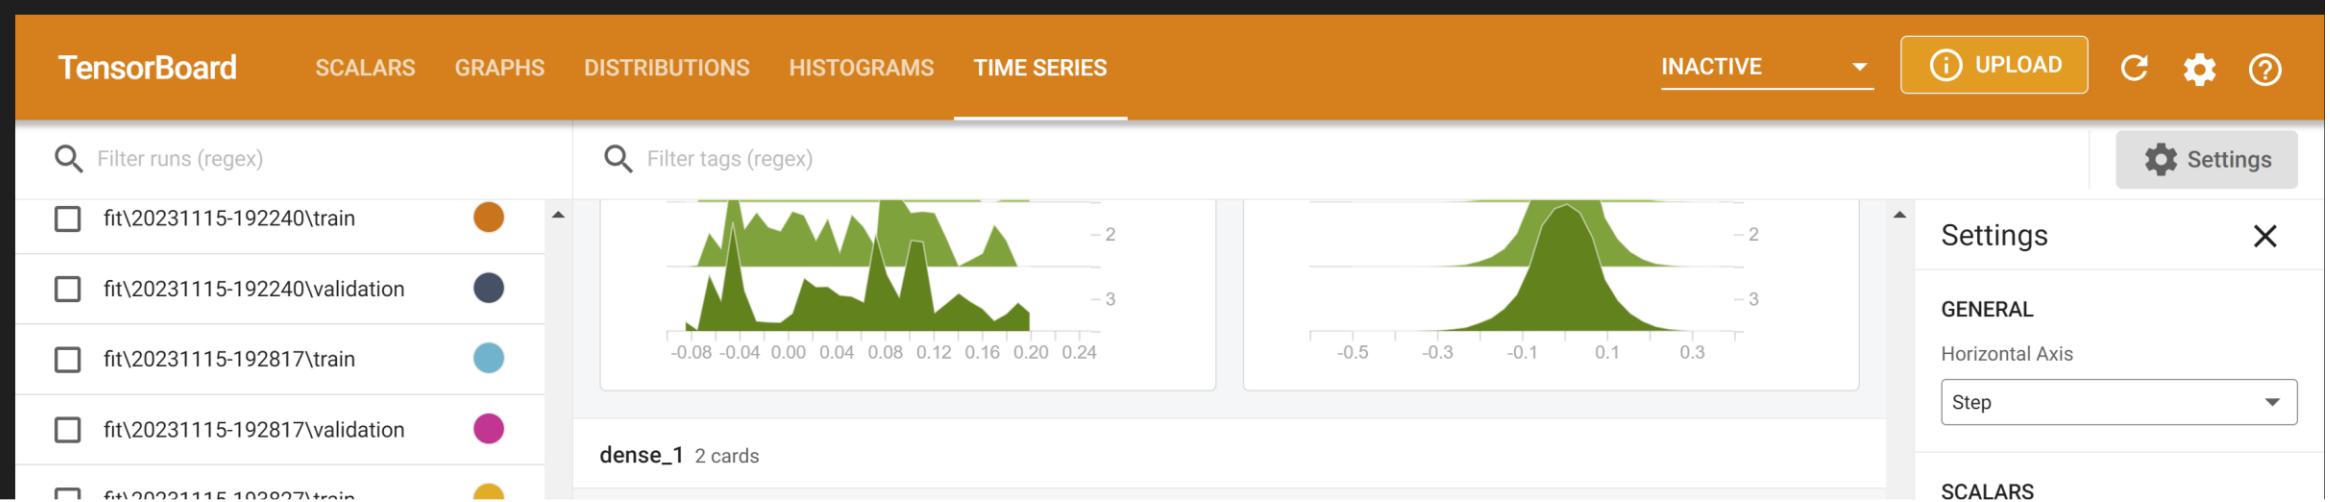
Python

Known TensorBoard instances:
- port 6006: logdir logs (started 0:00:06 ago; pid 31792)

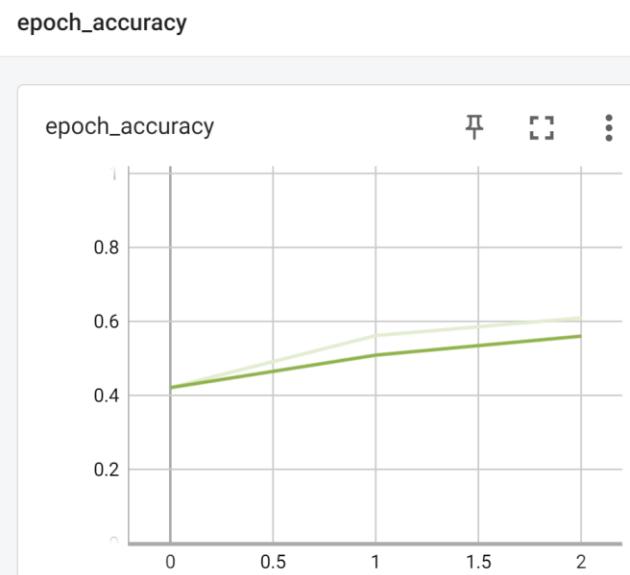
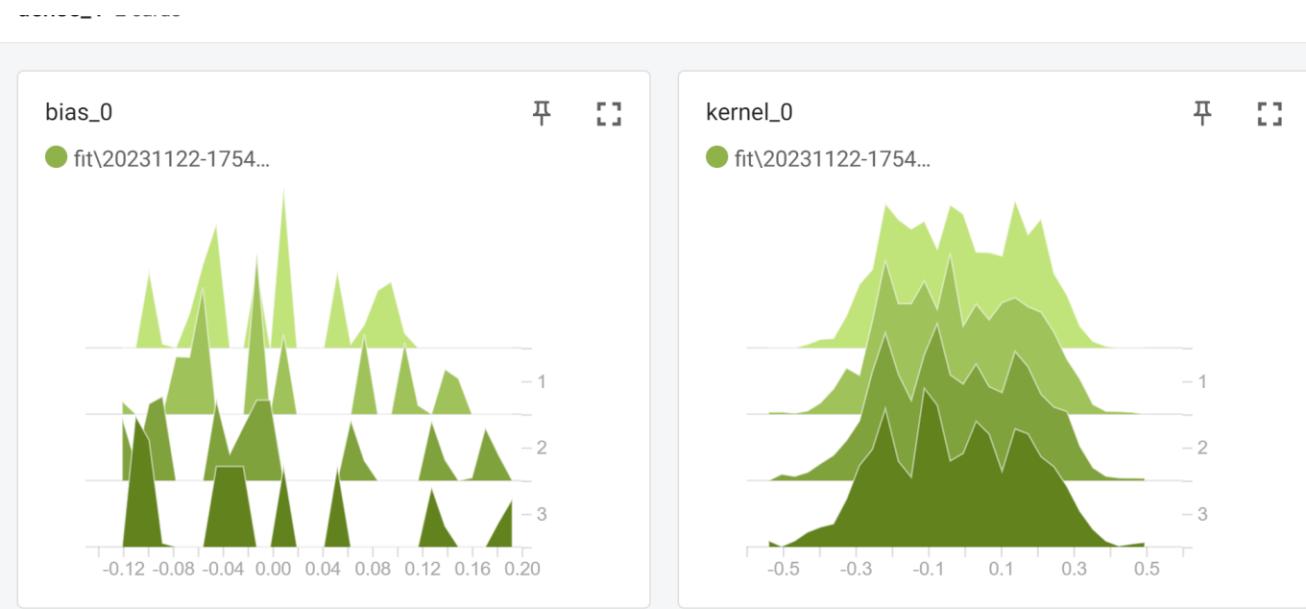
```
notebook.display(port=6006, height=1000)  
✓ 0.0s
```

Python

Selecting TensorBoard with logdir logs (started 0:00:12 ago; port 6006, pid 31792).



- fit\20231115-193827\train ●
- fit\20231115-193827\validation ●
- fit\20231115-194148\train ●
- fit\20231115-194148\validation ●
- fit\20231115-194424\train ●
- fit\20231115-194424\validation ●
- fit\20231115-203122\train ●
- fit\20231115-203122\validation ●
- fit\20231116-095942\train ●
- fit\20231116-095942\validation ●
- fit\20231116-100436\train ●
- fit\20231116-100436\validation ●
- fit\20231116-101113\train ●
- fit\20231116-101113\validation ●
- fit\20231116-101607\train ●
- fit\20231116-101607\validation ●
- fit\20231122-175455\train ●



SCALARS

Smoothing

Tooltip sorting method

Ignore outliers in chart scaling

Partition non-monotonic X axis ?

HISTOGRAMS

Mode

IMAGES

Brightness

Contrast

Show actual image size

```
# compile model
# optimizer: Adam
# loss: SparseCategoricalCrossentropy - when there are two or more label classes and labels are provided as integers
# metric - accuracy
# default batch_size = 32
model.compile(optimizer='adam',
               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
               metrics=['accuracy'])

log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Train model with 20 epochs
# Get loss/accuracy also on the test set
history = model.fit(train_images, train_labels, epochs=20,
                     validation_data=(test_images, test_labels),
                     callbacks=[tensorboard_callback])
```

```
Epoch 1/20
1563/1563 [=====] - 7s 4ms/step - loss: 1.5614 - accuracy: 0.4249 - val_loss: 1.2551 - val_accuracy: 0.5460
Epoch 2/20
1563/1563 [=====] - 6s 4ms/step - loss: 1.2293 - accuracy: 0.5608 - val_loss: 1.0759 - val_accuracy: 0.6217
Epoch 3/20
1563/1563 [=====] - 6s 4ms/step - loss: 1.0935 - accuracy: 0.6109 - val_loss: 1.0063 - val_accuracy: 0.6411
Epoch 4/20
1563/1563 [=====] - 6s 4ms/step - loss: 1.0132 - accuracy: 0.6400 - val_loss: 0.9667 - val_accuracy: 0.6531
Epoch 5/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.9518 - accuracy: 0.6634 - val_loss: 0.8780 - val_accuracy: 0.6941
Epoch 6/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.9098 - accuracy: 0.6786 - val_loss: 0.8656 - val_accuracy: 0.6974
Epoch 7/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.8693 - accuracy: 0.6924 - val_loss: 0.8335 - val_accuracy: 0.7038
Epoch 8/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.8440 - accuracy: 0.7004 - val_loss: 0.8259 - val_accuracy: 0.7086
Epoch 9/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.8105 - accuracy: 0.7158 - val_loss: 0.8046 - val_accuracy: 0.7189
Epoch 10/20
1563/1563 [=====] - 6s 4ms/step - loss: 0.7932 - accuracy: 0.7185 - val_loss: 0.8389 - val_accuracy: 0.7090
Epoch 11/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.7759 - accuracy: 0.7254 - val_loss: 0.8048 - val_accuracy: 0.7154
Epoch 12/20
1563/1563 [=====] - 10s 7ms/step - loss: 0.7575 - accuracy: 0.7309 - val_loss: 0.7812 - val_accuracy: 0.7263
Epoch 13/20
...
Epoch 19/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.6615 - accuracy: 0.7670 - val_loss: 0.7557 - val_accuracy: 0.7368
Epoch 20/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.6493 - accuracy: 0.7691 - val_loss: 0.7725 - val_accuracy: 0.7355
```

```
# Model accuracy

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'test_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')
```

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay

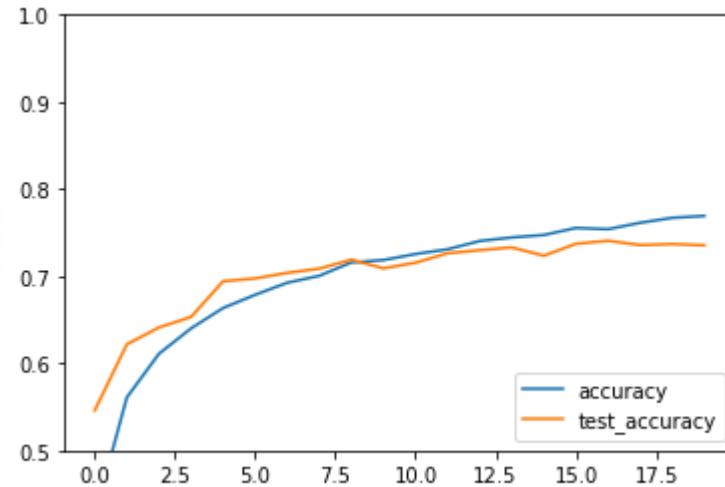
y_pred = model.predict(test_images)
predicted_categories = tf.argmax(y_pred, axis=1)

cm = confusion_matrix(test_labels, predicted_categories)

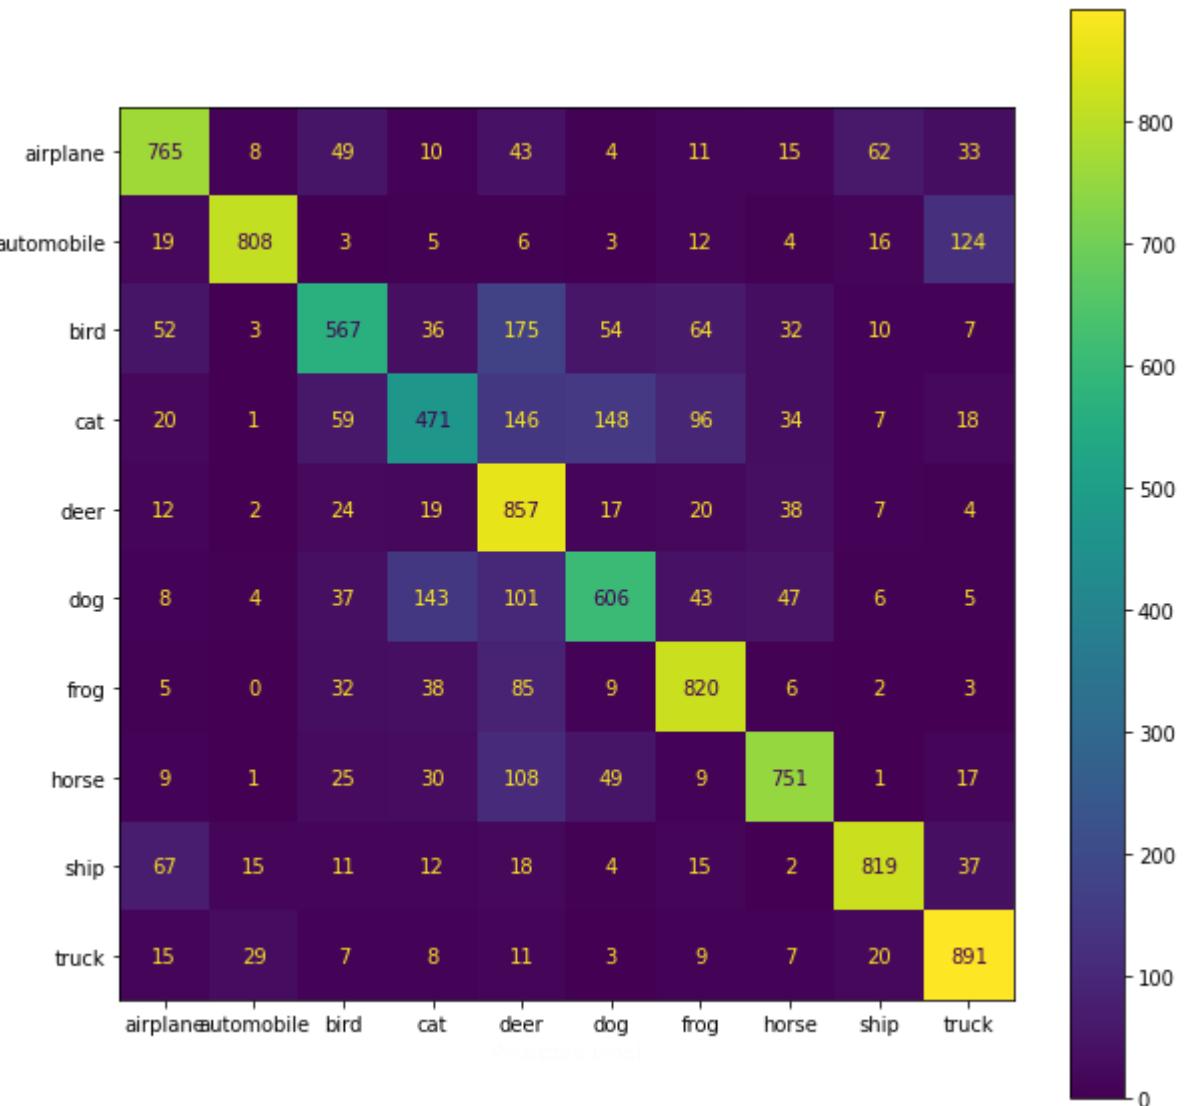
cmp = ConfusionMatrixDisplay(cm, display_labels = class_names)
fig, ax = plt.subplots(figsize=(10,10))
cmp.plot(ax = ax)
```

```
from sklearn.metrics import classification_report

print(classification_report(test_labels, predicted_categories, labels = list(range(10)), target_names= class_names))
```



	precision	recall	f1-score	support
airplane	0.79	0.77	0.78	1000
automobile	0.93	0.81	0.86	1000
bird	0.70	0.57	0.63	1000
cat	0.61	0.47	0.53	1000
deer	0.55	0.86	0.67	1000
dog	0.68	0.61	0.64	1000
frog	0.75	0.82	0.78	1000
horse	0.80	0.75	0.78	1000
ship	0.86	0.82	0.84	1000
truck	0.78	0.89	0.83	1000
accuracy			0.74	10000
macro avg	0.74	0.74	0.73	10000
weighted avg	0.74	0.74	0.73	10000



Prediction on test

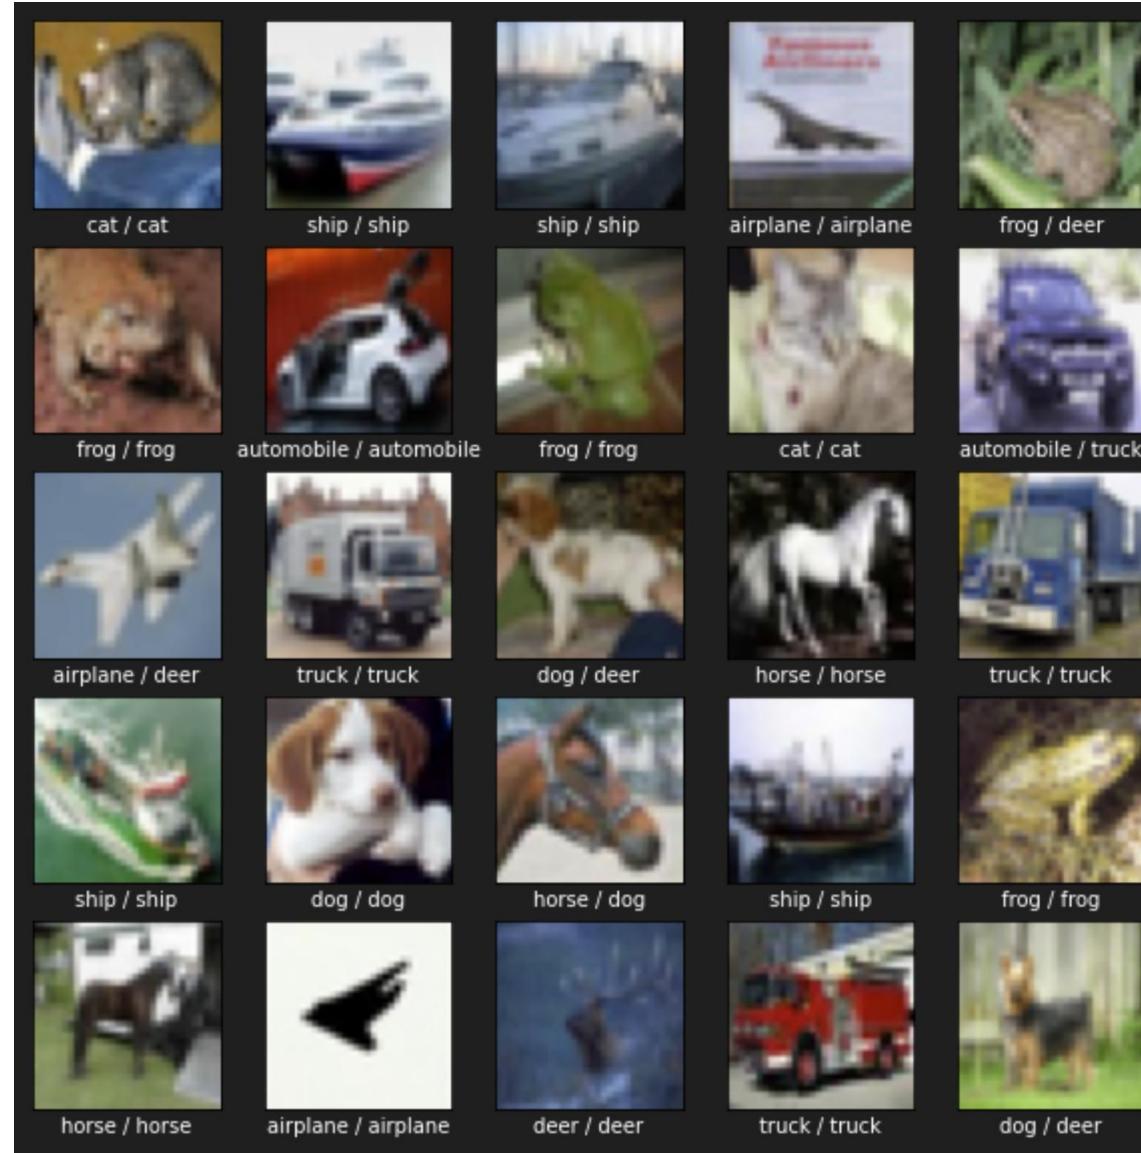
```
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(test_images[i])

    # Model was trained on batches,
    # hence the format expected for prediction is a tensor of the form
    # [batch_size, length, width, number of channels]
    # A picture means a batch_size = 1 => add a dimension to x
    x = np.expand_dims(test_images[i], axis=0)

predict_x = model.predict(x)

# The most probable class is taken from the resulting class probabilities
classes_x=np.argmax(predict_x,axis=-1)

# label showing real class/predicted class
plt.xlabel(class_names[test_labels[i][0]] + ' / ' + class_names[classes_x[0]])
plt.show()
```



CNN visualization

- Visual tutorial on CNN functioning
 - <https://poloclub.github.io/cnn-explainer/>
- Tool to visualize personal models
 - <https://github.com/zetane/viewer>



CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!



DeepExplainer

- Determines SHAP values by an improved version of the DeepLIFT algorithm
 - DeepLIFT backpropagates the neuron contributions to the input features
 - Contribution (positive or negative) = difference between neuron activation and reference activation
- The expectations are estimated on the base of selected background examples
 - 100 samples represent a good estimate; 1000 a very good one, but costly
- Interpretation:
 - Red pixels (positive) denote the features whose presence is significant for the class
 - Blue pixels (negative) show the features whose absence is significant for the class

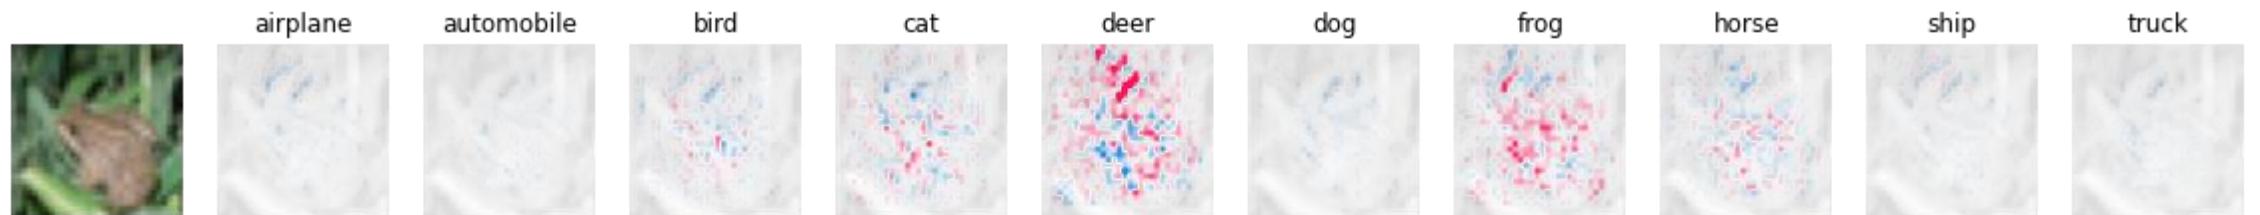
```
# DeepExplainer SHAP

import shap
import numpy as np

# select a background set of 100 examples to get the expectation
examples = train_images[np.random.choice(train_images.shape[0], 100, replace=False)]
explainer = shap.DeepExplainer(model, examples, class_names)

# explain model predictions on several test images
# red pixels increase the model output (presence of that area favors output)
# blue pixels decrease the output (absence of that area favors output)
for i in range(25):
    print(class_names[test_labels[i][0]])
    shap_values = explainer.shap_values(test_images[i:i+1])
    shap.image_plot(shap_values, test_images[i:i+1], labels=class_names)
```





-0.0075 -0.0050 -0.0025 0.0000 0.0025 0.0050 0.0075

label index



-0.020 -0.015 -0.010 -0.005 0.000 0.005 0.010 0.015 0.020

label index

Transfer learning

- Deep learning needs
 - Big data
 - Big computing resources
- Take the parameters from an already trained network on a large data set
 - Data: ImageNet, CIFAR
 - Pre-trained models: VGG, Inception, AlexNet, ResNet
 - Initial layers have learnt general features
- Train the final layers for the problem at hand
 - Learn the specific features of the current data
- Hence
 - Resolve problem with scarce available data
 - Less parameters to train

VGG-19

- Common choice for transfer learning
- 19 layers (16 convolutional + 3 feed-forward)
- Top layers are the last layers added to the network

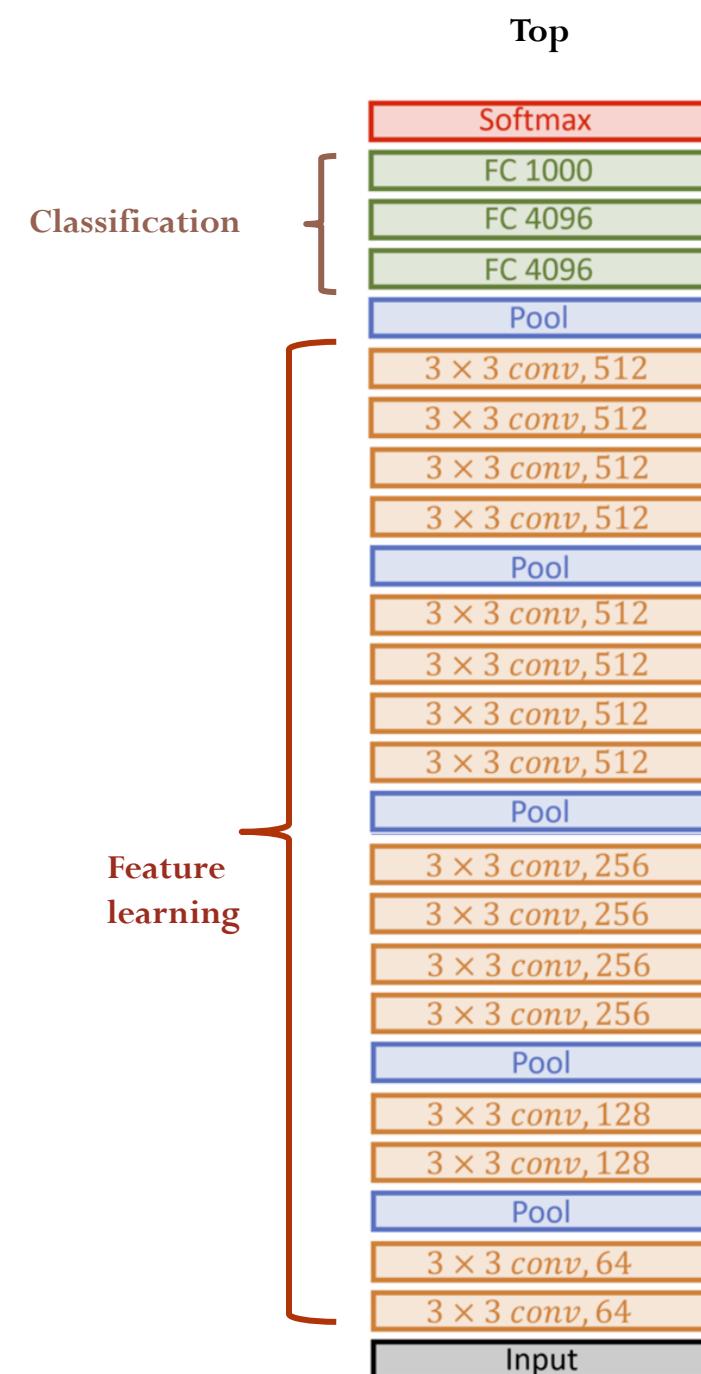
```
# Transfer Learning with VGG19

import tensorflow as tf
from tensorflow.keras import datasets, layers, models
from keras.applications.vgg19 import VGG19
import matplotlib.pyplot as plt
import numpy as np

# load CIFAR10 data set

(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# normalize pixel values
train_images, test_images = train_images / 255.0, test_images / 255.0
```



```
# Base: VGG19 pretrained on the ImageNet data set -> set of weights of the pretraining
# include_top: (not) include the 3 fully-connected layers (those for classification) at the top of the network
# i.e., take just the feature extraction convolutional layers

TLmodel = VGG19(include_top = False, weights = 'imagenet', classes= train_labels)

# keep all layers as trained on ImageNet (i.e. frozen), except the last given number
# i.e., the first layer have learned the basic features
# the last ones will learn the characteristics specific to the current CIFAR problem

for layer in TLmodel.layers[:-6]: # just the last given number of layers are trainable
    layer.trainable = False

# see all trainable layers
for layer in TLmodel.layers:
    print(layer, layer.trainable)

# create a sequential model that begins with the pretrained VGG + the last trainable layers
model = models.Sequential()

# increase image size to be better fitted with the VGG19
model.add(layers.UpSampling2D())
model.add(layers.UpSampling2D())

# add the VGG19
model.add(TLmodel)

# The flatten layer collapses the spatial dimensions of the input into 1D, and then fully connected follows
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.7))
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.3))

# Add the final dense layer with the output for the 10 classes
model.add(layers.Dense(10, activation='softmax'))
```

```
<keras.engine.input_layer.InputLayer object at 0x0000023496ADD730> False
<keras.layers.convolutional.Conv2D object at 0x0000023496ADD3D0> False
<keras.layers.convolutional.Conv2D object at 0x0000023496ADD880> False
<keras.layers.pooling.MaxPooling2D object at 0x000002348E35AB80> False
<keras.layers.convolutional.Conv2D object at 0x000002348E321FD0> False
<keras.layers.convolutional.Conv2D object at 0x0000023496AE1640> False
<keras.layers.pooling.MaxPooling2D object at 0x0000023496AB01C0> False
<keras.layers.convolutional.Conv2D object at 0x0000023496AF0340> False
<keras.layers.convolutional.Conv2D object at 0x000002348E5529A0> False
<keras.layers.convolutional.Conv2D object at 0x0000023496AF45B0> False
<keras.layers.convolutional.Conv2D object at 0x0000023497AD2BB0> False
<keras.layers.pooling.MaxPooling2D object at 0x0000023497AD2310> False
<keras.layers.convolutional.Conv2D object at 0x0000023497AD2520> False
<keras.layers.convolutional.Conv2D object at 0x0000023496AFDD90> False
<keras.layers.convolutional.Conv2D object at 0x0000023497ADED00> False
<keras.layers.convolutional.Conv2D object at 0x0000023497ADE0D0> False
<keras.layers.pooling.MaxPooling2D object at 0x0000023497AE60D0> True
<keras.layers.convolutional.Conv2D object at 0x0000023496AF4F70> True
<keras.layers.convolutional.Conv2D object at 0x0000023497ADEBB0> True
<keras.layers.convolutional.Conv2D object at 0x0000023497ADEF70> True
<keras.layers.convolutional.Conv2D object at 0x0000023497AF2AC0> True
<keras.layers.pooling.MaxPooling2D object at 0x0000023497AFE9D0> True
```

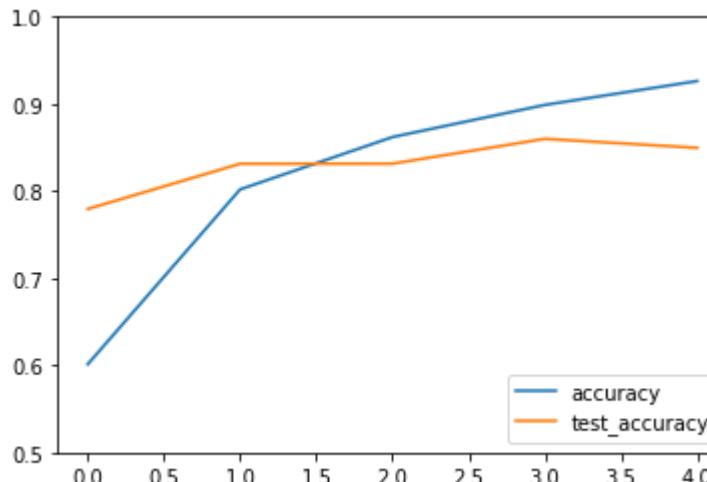
```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),
              loss=tf.keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
```

```
history = model.fit(train_images, train_labels, epochs=5,
                      validation_data=(test_images, test_labels))
```

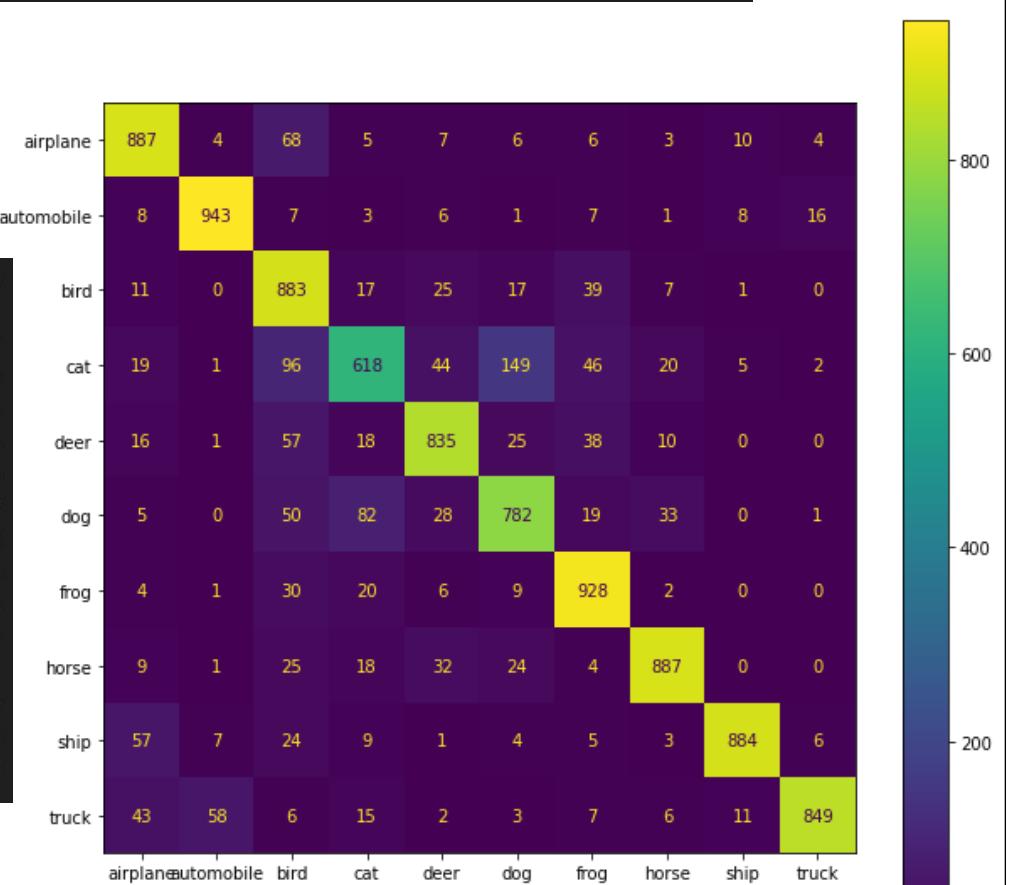
```

Epoch 1/5
1563/1563 [=====] - 115s 73ms/step - loss: 1.1800 - accuracy: 0.6012 - val_loss: 0.6623 - val_accuracy: 0.7794
Epoch 2/5
1563/1563 [=====] - 147s 94ms/step - loss: 0.6321 - accuracy: 0.8019 - val_loss: 0.5281 - val_accuracy: 0.8313
Epoch 3/5
1563/1563 [=====] - 176s 113ms/step - loss: 0.4448 - accuracy: 0.8619 - val_loss: 0.5440 - val_accuracy: 0.8314
Epoch 4/5
1563/1563 [=====] - 146s 94ms/step - loss: 0.3236 - accuracy: 0.8988 - val_loss: 0.4772 - val_accuracy: 0.8599
Epoch 5/5
1563/1563 [=====] - 158s 101ms/step - loss: 0.2378 - accuracy: 0.9262 - val_loss: 0.5607 - val_accuracy: 0.8496

```



	precision	recall	f1-score	support
airplane	0.91	0.81	0.86	1000
automobile	0.87	0.94	0.90	1000
bird	0.88	0.73	0.80	1000
cat	0.67	0.70	0.68	1000
deer	0.77	0.84	0.80	1000
dog	0.77	0.72	0.75	1000
frog	0.80	0.92	0.86	1000
horse	0.89	0.85	0.87	1000
ship	0.89	0.93	0.91	1000
truck	0.90	0.89	0.90	1000
accuracy			0.83	10000
macro avg	0.84	0.83	0.83	10000
weighted avg	0.84	0.83	0.83	10000



HOW TO CONFUSE MACHINE LEARNING

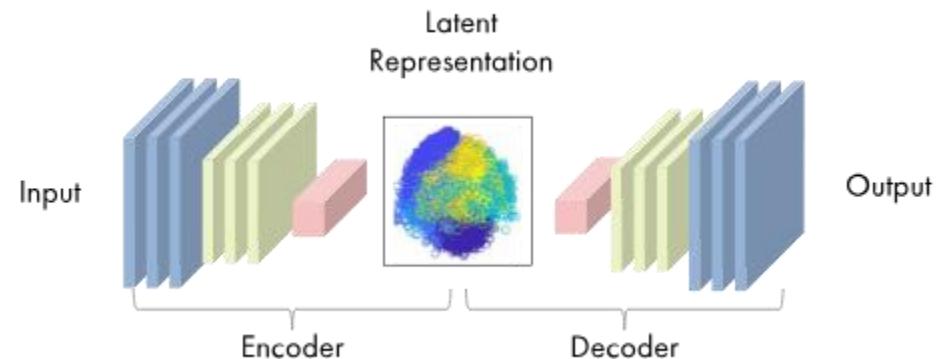


Balbix®

Autoencoders

Autoencoders

- Unsupervised learning
- Used for data compression



mathworks.com

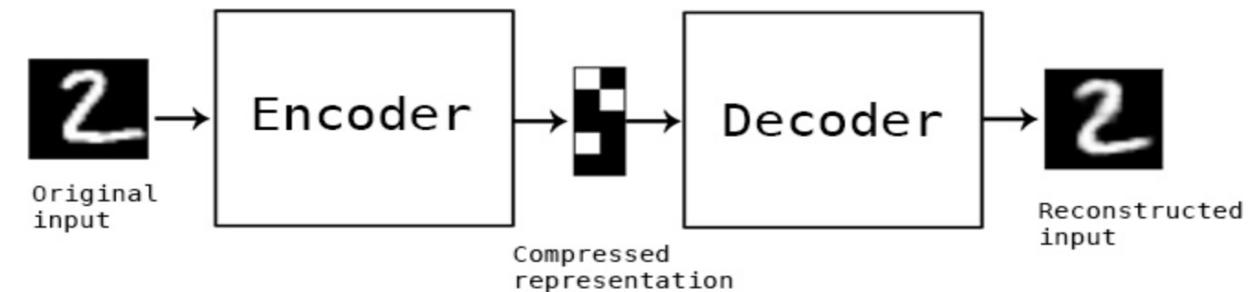
- 2 stages
 - Encoder: Reduces dimensionality, learns the most important features
 - Decoder: Reconstructs the image from the reduced representation
- Minimize reconstruction loss
- Dimensionality reduction
- Data denoising

Examples

- Convolutional autoencoder
 - Convolutional layers instead of dense
- Sequence-to-sequence autoencoder
 - Recurrent network layers (LSTM)
- Variational autoencoder
 - It learns a probability distribution model of the data
 - New points can be sampled from this distribution -> a type of generative model

Autoencoders for anomaly detection

- Train the autoencoder to learn normal data samples
- Weights will be obtained such that reconstruction error is minimized
- At inference, any abnormal data will get a very high reconstruction error
- A threshold can be appointed to signal anomalies



```
import numpy as np
import matplotlib.pyplot as plt
import keras
from sklearn.model_selection import train_test_split

(x_train, _), (x_test, _) = keras.datasets.mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((-1, 28 * 28))
x_test = x_test.reshape((-1, 28 * 28))

# the autoencoder
model = keras.Sequential([
    # encoder
    keras.layers.Dense(128, activation='relu', input_shape=(x_train.shape[1],)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(32, activation='relu'),

    # decoder
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(x_train.shape[1], activation='sigmoid') # ensures reconstructed pixel value between 0 (black) and 1 (white)
])

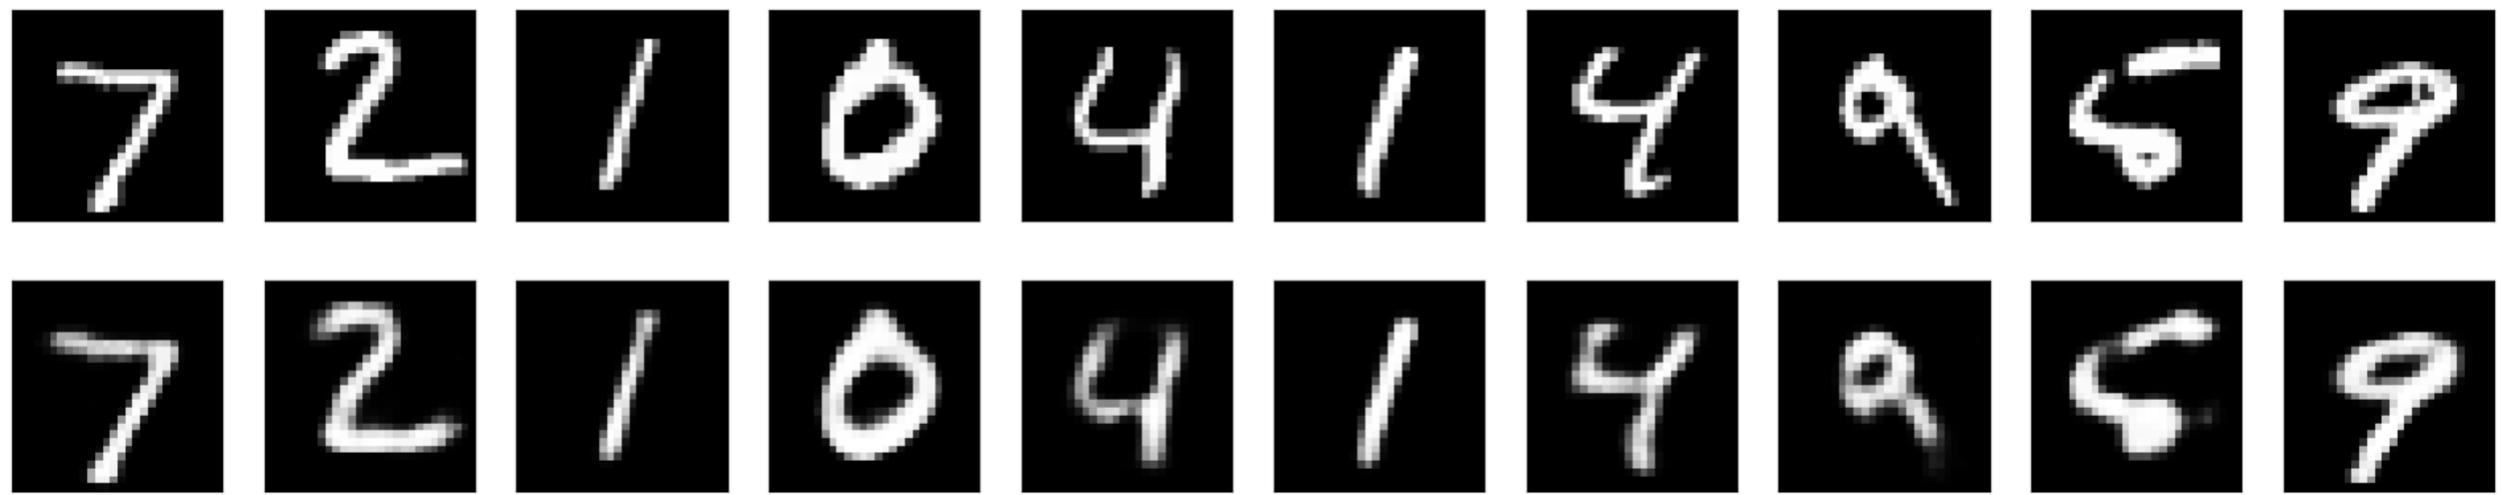
model.compile(optimizer='adam', loss='binary_crossentropy') # per-pixel binary crossentropy loss

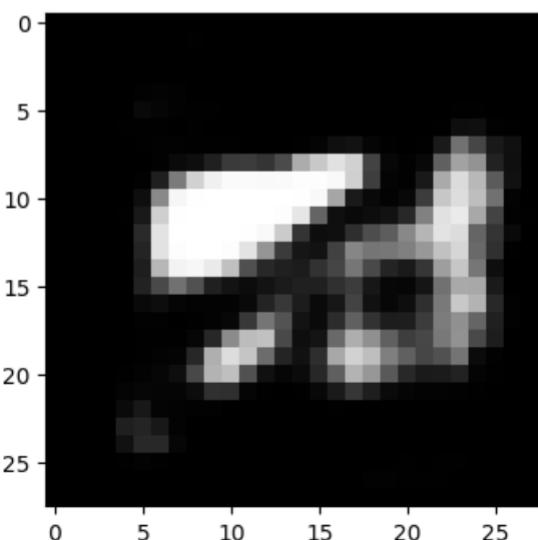
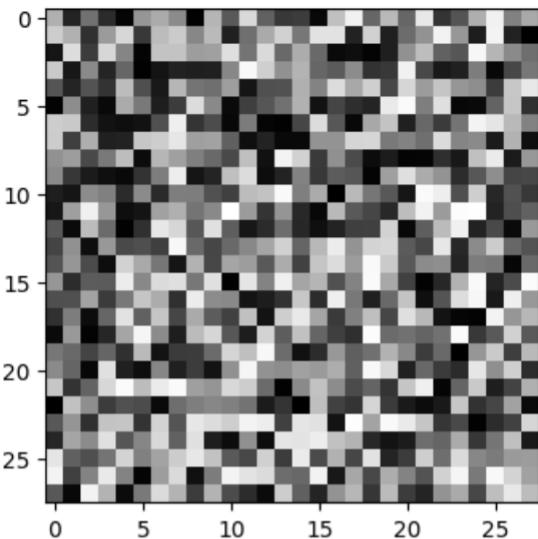
# Train the model on the usual MNIST data
history = model.fit(x_train, x_train, epochs=20, batch_size=256, validation_data=(x_test, x_test))
```

```
# Calculate reconstruction loss for the normal images
reconstruction_normal = model.predict(x_test)
reconstruction_loss_normal = np.mean(np.abs(x_test - reconstruction_normal), axis=1)

# Print average reconstruction loss
print(f"Average Reconstruction Loss for Normal Data: {np.mean(reconstruction_loss_normal)}")
```

313/313 ————— 1s 2ms/step
Average Reconstruction Loss for Normal Data: 0.03286447748541832





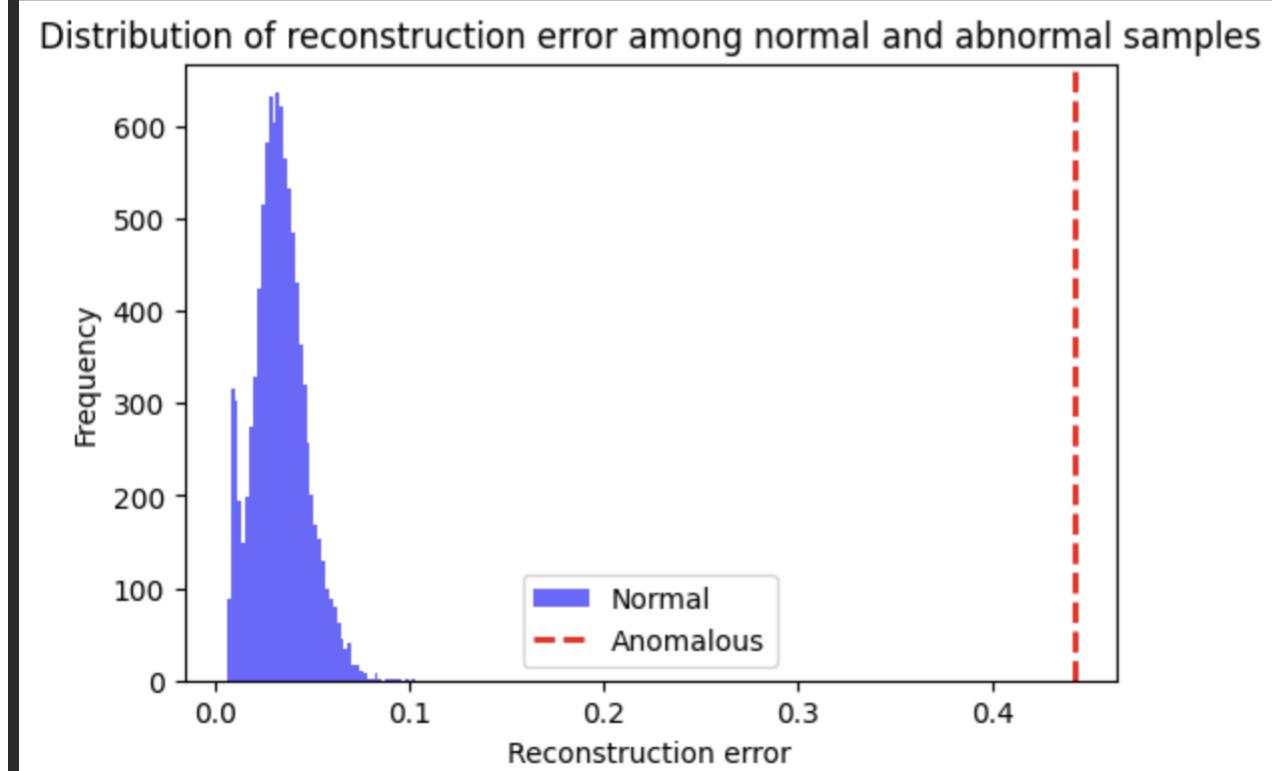
```
# Detect an anomaly following reconstruction error
# This could be an untypical image, for example a random 28x28 image
anomalous_image = np.random.rand(28 * 28)

# Calculate reconstruction loss for the anomalous images
reconstruction_anomalous = model.predict(np.array([anomalous_image]))
reconstruction_loss_anomalous = np.mean(np.abs(np.array([anomalous_image]) - reconstruction_anomalous), axis=1)

print(f'Reconstruction Loss for Anomalous Data: {reconstruction_loss_anomalous[0]}')

1/1 ━━━━━━━━ 0s 45ms/step
Reconstruction Loss for Anomalous Data: 0.44151989400984815
```

```
# Visualization of the distribution of the reconstruction error
plt.figure(figsize=(6, 4))
plt.hist(reconstruction_loss_normal, bins=50, alpha=0.6, color='b', label='Normal')
plt.axvline(x=reconstruction_loss_anomalous[0], color='r', linestyle='dashed', linewidth=2, label='Anomalous')
plt.title('Distribution of reconstruction error among normal and abnormal samples')
plt.xlabel('Reconstruction error')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



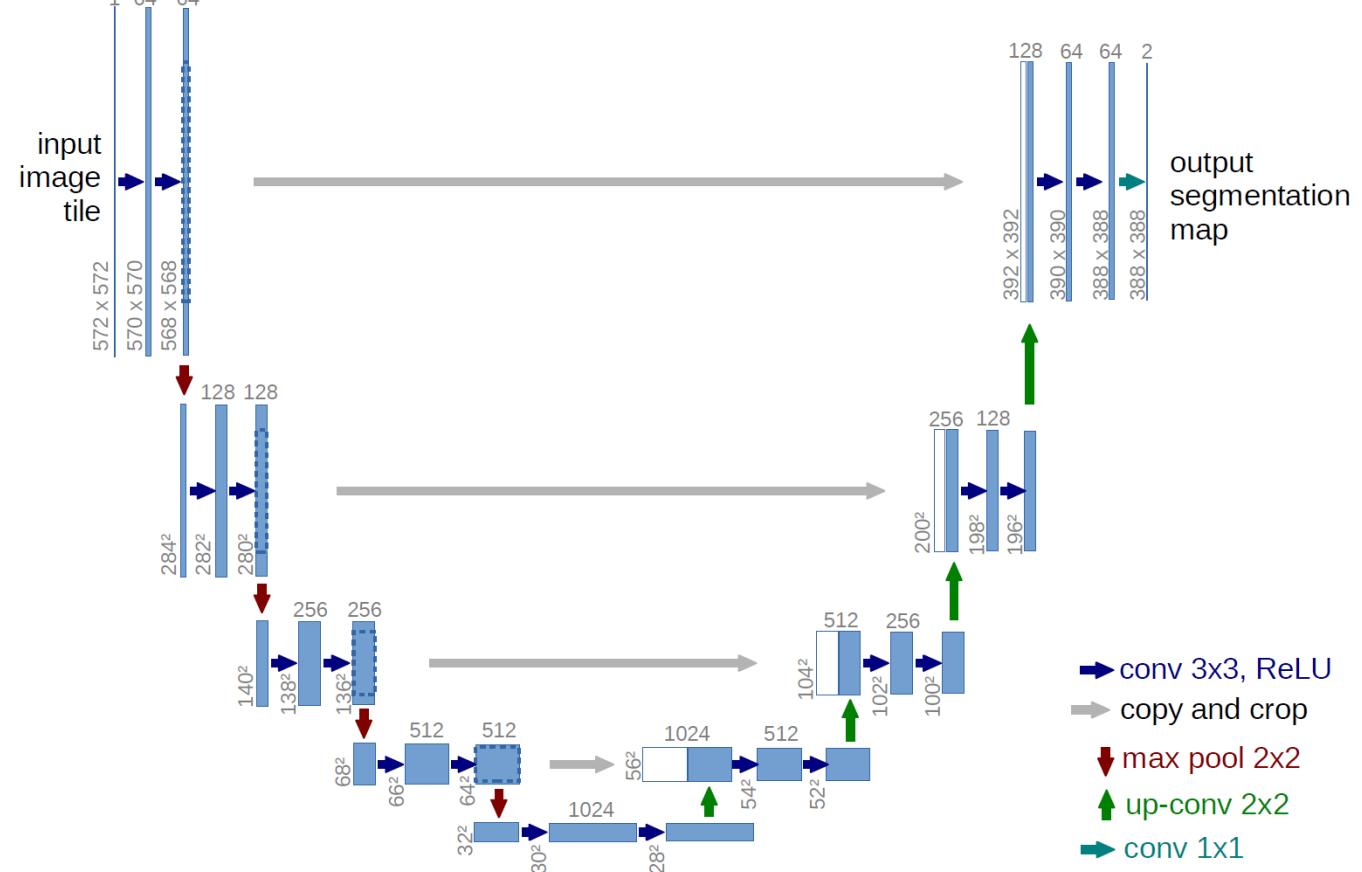
Semantic segmentation

Competing architectures

U-Net	Mask R-CNN
Fully CNN	Extension of Faster Regional CNN
Encoder-decoder architecture	Object detection in found region proposals
Skip connections for direct transfer of features from encoder to decoder	Mask prediction in parallel to detection through a CNN
Designed for fine-grained detailed biomedical segmentation	General purpose image segmentation

U-Net

- Addresses pixel-wise representations of the images
 - by downsampling (encoder)
 - and upsampling (decoder)
- Skip connections – send information at the same size

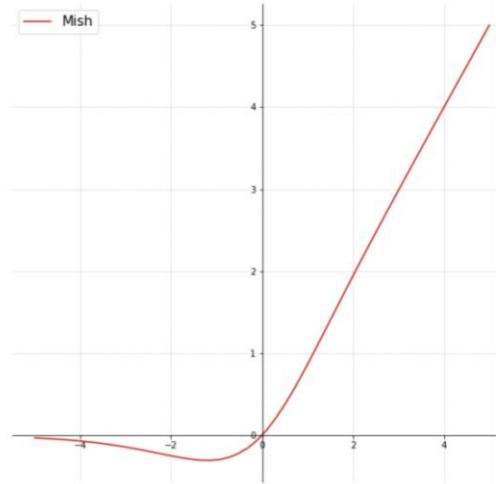


U-Net with ResNet34 on CamVid in Pytorch

- CamVid – frames from video of cars and pedestrians
- <https://mi.eng.cam.ac.uk/research/projects/VideoRec/CamVid/>
- Batch normalization – normalizing and rescaling
 - follows convolutional layers and lies before the ReLU layers
 - allows higher learning rates and reduces the strong dependency on initialization
- Self-attention layers – to suppress activations in irrelevant regions, i.e. features carried across by the skip connections, and highlight the relevant zones.
- Pixel accuracy as performance metric
- Mish activation function, Ranger optimizer
- First time for the fastai deep learning library
 - !pip install fastai wwf -q --upgrade



```
from fastai.vision.all import *
from fastcore.xtras import Path
from fastai.callback.hook import summary
from fastai.callback.progress import ProgressCallback
from fastai.data.block import DataBlock
from fastai.data.external import untar_data, URLs
from fastai.data.transforms import get_image_files, FuncSplitter, Normalize
from fastai.layers import Mish
from fastai.losses import BaseLoss
from fastai.optimizer import ranger
from fastai.torch_core import tensor
from fastai.vision.augment import aug_transforms
from fastai.vision.core import PILImage, PILMask
from fastai.vision.data import ImageBlock, MaskBlock, imagenet_stats
from fastai.vision.learner import unet_learner
from PIL import Image
import numpy as np
from torch import nn
from torchvision.models.resnet import resnet34
import torch
import torch.nn.functional as F
```



```
▶ # data set CamVid (The Cambridge-driving Labeled Video Database)
# 701 frames, 32 classes manually segmented
path = untar_data(URLs.CAMVID)
```

```
→ [ ] 100.00% [598917120/598913237 00:16<00:00]
```

```
[ ] # get validation set
valid_fnames = (path/'valid.txt').read_text().split('\n')

def FileSplitter(fname):
    valid = Path(fname).read_text().split('\n')
    def _func(x): return x.name in valid
    def _inner(o, **kwargs): return FuncSplitter(_func)(o)
    return _inner
```

```
 # view samples of the data set  
path_im = path/'images'  
path_lbl = path/'labels'  
fnames = get_image_files(path_im)  
lbl_names = get_image_files(path_lbl)  
img_fn = fnames[20]  
img = PILImage.create(img_fn)  
img.show(figsize=(5,5))
```

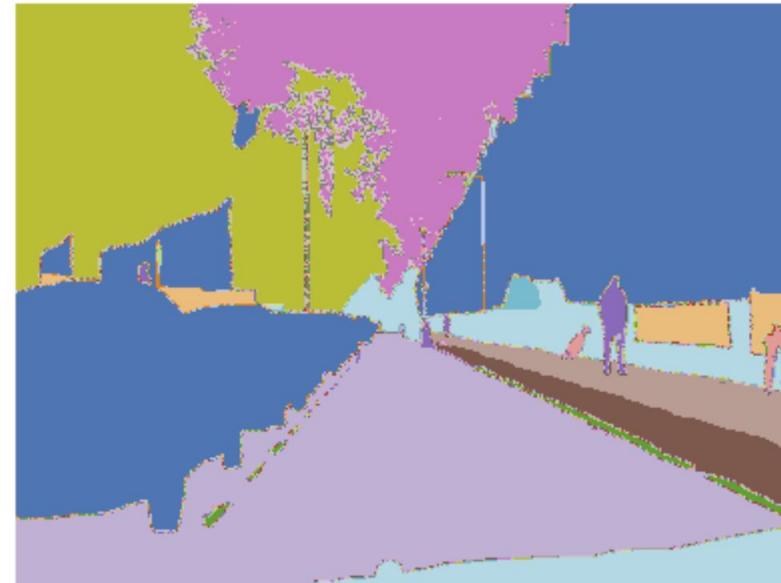


<Axes: >



```
▶ # see mask  
get_msk = lambda o: path/'labels'/f'{o.stem}_P{o.suffix}'  
msk = PILMask.create(get_msk(img_fn))  
msk.show(figsize=(5,5), alpha=1)
```

→ <Axes: >



```
[ ] # see the inside values for the mask  
tensor(msk)
```

```
tensor([[26, 26, 26, ..., 4, 4, 4],  
        [26, 26, 26, ..., 4, 4, 4],  
        [26, 26, 26, ..., 4, 4, 4],  
        ...,  
        [17, 17, 17, ..., 30, 30, 30],  
        [17, 17, 17, ..., 30, 30, 30],  
        [17, 17, 17, ..., 30, 30, 30]], dtype=torch.uint8)
```

```
▶ # see classes of objects in the images

codes = np.loadtxt(path/'codes.txt', dtype=str); codes

[ ] array(['Animal', 'Archway', 'Bicyclist', 'Bridge', 'Building', 'Car',
       'CartLuggagePram', 'Child', 'Column_Pole', 'Fence', 'LaneMkgsDriv',
       'LaneMkgsNonDriv', 'Misc_Text', 'MotorcycleScooter', 'OtherMoving',
       'ParkingBlock', 'Pedestrian', 'Road', 'RoadShoulder', 'Sidewalk',
       'SignSymbol', 'Sky', 'SUVPickupTruck', 'TrafficCone',
       'TrafficLight', 'Train', 'Tree', 'Truck_Bus', 'Tunnel',
       'VegetationMisc', 'Void', 'Wall'], dtype='<U17')
```

```
[ ] # make a dictionary
name2id = {v:k for k,v in enumerate(codes)}
void_code = name2id['Void']
```

```
[ ] # make a data loader
camvid = DataBlock(blocks=(ImageBlock, MaskBlock(codes)),
                  get_items=get_image_files,
                  splitter=FileSplitter(path/'valid.txt'),
                  get_y=get_msk,
                  batch_tfms=[*aug_transforms(), Normalize.from_stats(*imagenet_stats)])
```

```
[ ] # batch size equal to 1
dls = camvid.dataloaders(path/"images", bs=1)
```

```
[ ] # accuracy function

# output values are squeezed in a segmentation mask matrix
# for each pixel compare argmax to the target mask
# get average
def acc_camvid(inp, targ):
    targ = targ.squeeze(1)
    mask = targ != void_code
    return (inp.argmax(dim=1)[mask]==targ[mask]).float().mean()
```

▶ # the U-Net model

```
opt = ranger # the optimizer
lr = 1e-3
dls.vocab = codes

# add self-attention layers
# Mish activation function
learn = unet_learner(dls, resnet34, metrics=acc_camvid, self_attention=True, act_cls=Mish, opt_func=opt)
```

→ Downloading: "<https://download.pytorch.org/models/resnet34-b627a593.pth>" to /root/.cache/torch/hub/checkpoints/resnet34-b627a593.pth
100%|██████████| 83.3M/83.3M [00:01<00:00, 70.9MB/s]

```
[ ] learn.summary()
```

DynamicUnet (Input shape: 1 x 3 x 720 x 960)			
Layer (type)	Output Shape	Param #	Trainable
	1 x 64 x 360 x 480		
Conv2d		9408	False
BatchNorm2d		128	True
ReLU			
	1 x 64 x 180 x 240		
MaxPool2d			
Conv2d		36864	False
BatchNorm2d		128	True
ReLU			
Conv2d		36864	False
BatchNorm2d		128	True
Conv2d		36864	False
BatchNorm2d		128	True
ReLU			
Conv2d		36864	False
BatchNorm2d		128	True
Conv2d		36864	False
BatchNorm2d		128	True
ReLU			
Conv2d		36864	False
BatchNorm2d		128	True
	1 x 384 x 360 x 480		
Conv2d		37248	True
Mish			
	1 x 96 x 720 x 960		
PixelShuffle			
ResizeToOrig			
	1 x 99 x 720 x 960		
MergeLayer			
Conv2d		88308	True
Mish			
Conv2d		88308	True
Sequential			
Mish			
	1 x 32 x 720 x 960		
Conv2d		3200	True
ToTensorBase			

```
▶ # fit the model  
  
# Fit-Flat-Cosine policy for fast training (superconvergence)  
# fit the model for the given number of epochs at flat learning rate before cosine annealing  
  
# cosine annealing - learning rate scheduler starting with the given maximum value  
# that is decreased fastly to the mininum and increased rapidly again  
# It is a learning restart and re-usage of good weights  
  
learn.fit_flat_cos(10, slice(lr))
```

epoch	train_loss	valid_loss	acc_camvid	time
0	0.706903	0.544879	0.852193	07:26
1	0.526745	0.426549	0.888225	07:33
2	0.469289	0.429322	0.893548	07:32
3	0.385277	0.355116	0.901144	07:34
4	0.368626	0.337369	0.911236	07:33
5	0.326502	0.294379	0.920019	07:32
6	0.338418	0.339668	0.907746	07:31
7	0.319459	0.318426	0.916108	07:31
8	0.251843	0.255840	0.927702	07:31
9	0.232762	0.279978	0.922318	07:30

```
[ ] learn.save('model')
```

```
learn.show_results(figsize=(18,8))
```



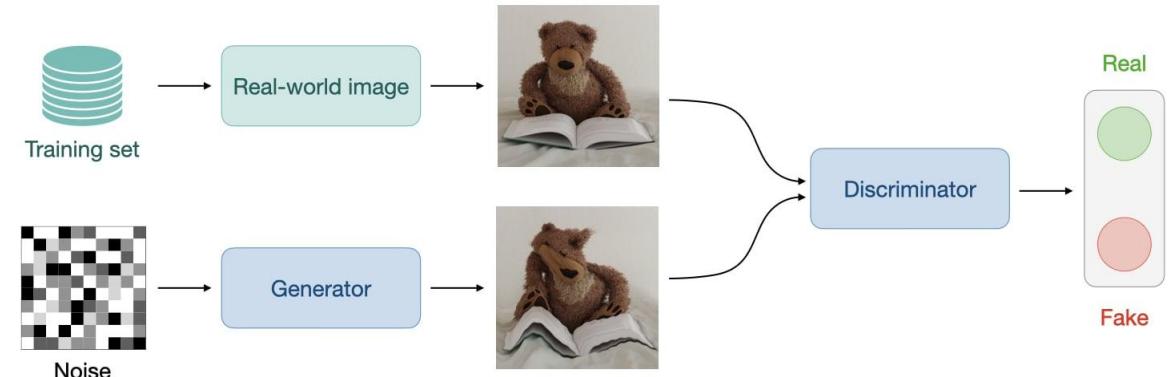
Target/Prediction



Generative AI

GenAI

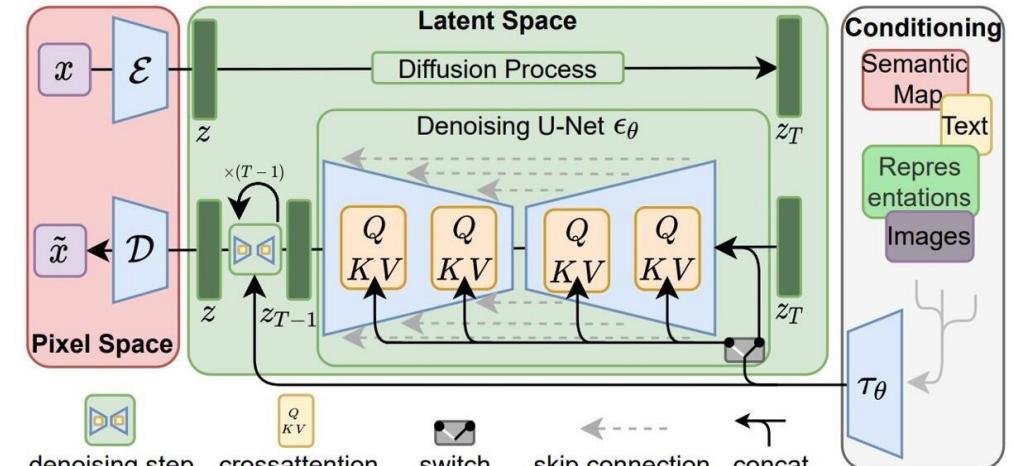
1. Learn from a very big number of examples
 2. Add inductive biases to insert new content based on context
- Generative adversarial networks
 - CoModGAN
 - Fast Fourier convolution
 - LaMa
 - Diffusion
 - Stable Diffusion



Stable Diffusion

1. Diffusion

- a. Forward: Infer Gaussian noise gradually
- b. Reverse: Conditional denoising
- c. U-Net for image2image mapping
- d. Modulate with condition of the text prompt



<https://arxiv.org/pdf/2112.10752>

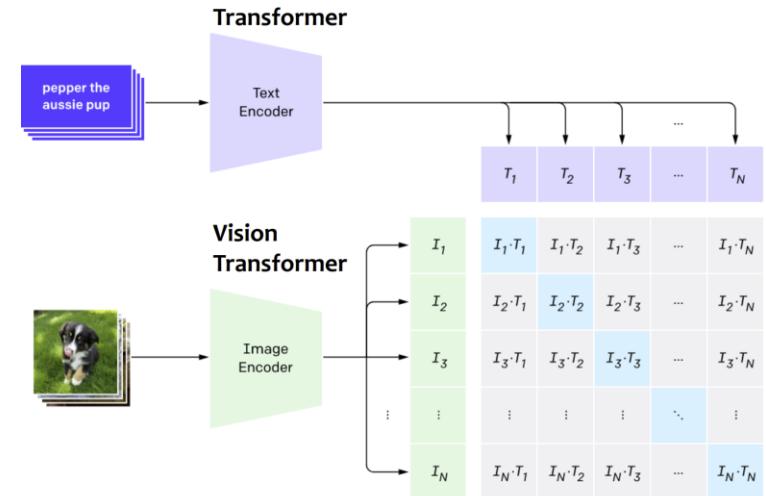
2. U-Net + **cross**-attention (combines two different input sequences)

3. Contrastive learning for prompts – CLIP model

- a. Associate image with name in a joint encoding
- b. Vision transformer + causal language model

4. Autoencoder

- a. Compresses images to increase training/generation speed
- b. Latent diffusion: it takes place in the latent space



<https://openai.com/research/clip>

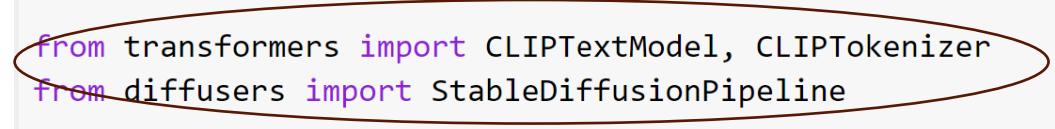
Stable Diffusion in Pytorch with Hugging Face

- Hugging Face – the machine learning platform known for Transformers and Diffusers libraries
 - The models hosted on Hugging Face can be imported in the workflow and creation of deep learning pipelines.
- <https://huggingface.co/>
- !pip install transformers diffusers lpips accelerate
- !jupyter nbextension enable --py widgetsnbextension
- from huggingface_hub import notebook_login
- notebook_login()
- !huggingface-cli whoami ## testing being logged in

Imports

```
import numpy
from tqdm.auto import tqdm

import torch
from torch import autocast
from torchvision import transforms as tfms

  
from transformers import CLIPTextModel, CLIPTokenizer  
from diffusers import StableDiffusionPipeline

from PIL import Image
from matplotlib import pyplot as plt

# For video display
from IPython.display import HTML
from base64 import b64encode

# For interactive notebook
import ipywidgets as widgets
from ipywidgets import interact, interact_manual

# Set device - setting the runtime to GPU
torch_device = "cuda" if torch.cuda.is_available() else "cpu"
print(torch_device)
```

Load pretrained SD

```
StableDiffusionPipeline {  
    "_class_name": "StableDiffusionPipeline",  
    "_diffusers_version": "0.23.1",  
    "_name_or_path": "CompVis/stable-diffusion-v1-4",  
    "feature_extractor": [  
        "transformers",  
        "CLIPImageProcessor"  
    ],  
    "requires_safety_checker": true,  
    "safety_checker": [  
        "stable_diffusion",  
        "StableDiffusionSafetyChecker"  
    ],  
    "scheduler": [  
        "diffusers",  
        "PNDMScheduler"  
    ],  
    "text_encoder": [  
        "transformers",  
        "CLIPTextModel"  
    ],  
    "tokenizer": [  
        "transformers",  
        "CLIPTokenizer"  
    ],  
    "unet": [  
        "diffusers",  
        "UNet2DConditionModel"  
    ],  
    "vae": [  
        "diffusers",  
        "AutoencoderKL"  
    ]  
}
```

```
stable_diff_pipe = StableDiffusionPipeline.from_pretrained('CompVis/stable-diffusion-v1-4')  
stable_diff_pipe.to(torch_device)
```

Prompt-to-image

```
prompt = ['A girl visiting Timisoara']

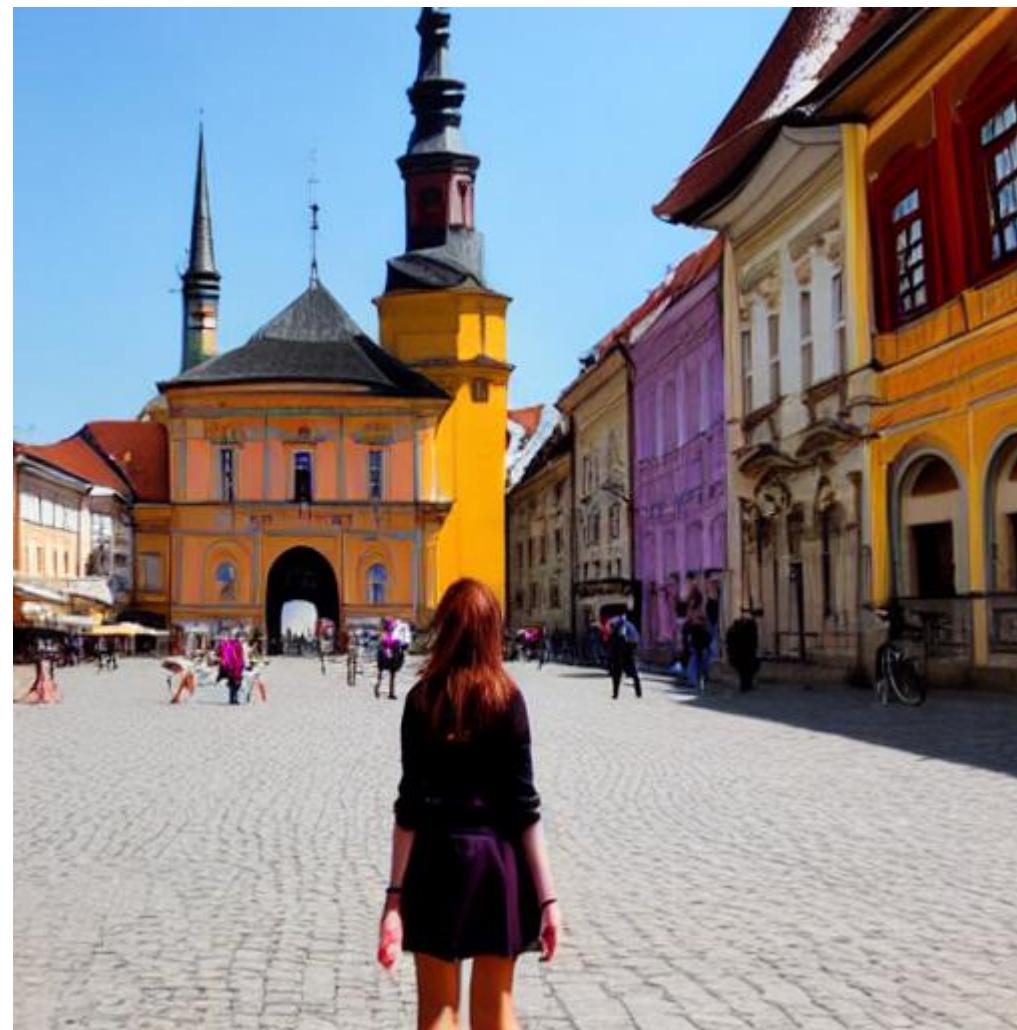
# the degree by which the image generation process follows the text prompt
# the higher the value, the more the image is related to the given prompt
guidance_scale = 12.5

# Change the seed parameter to create a different image each time
generator = torch.manual_seed(69)

with autocast("cuda"):
    image = stable_diff_pipe(prompt, guidance_scale = guidance_scale).images[0]

image
```

“A girl visiting Timisoara” result



Semantic Inpainting

- !pip install --pre -U xformers

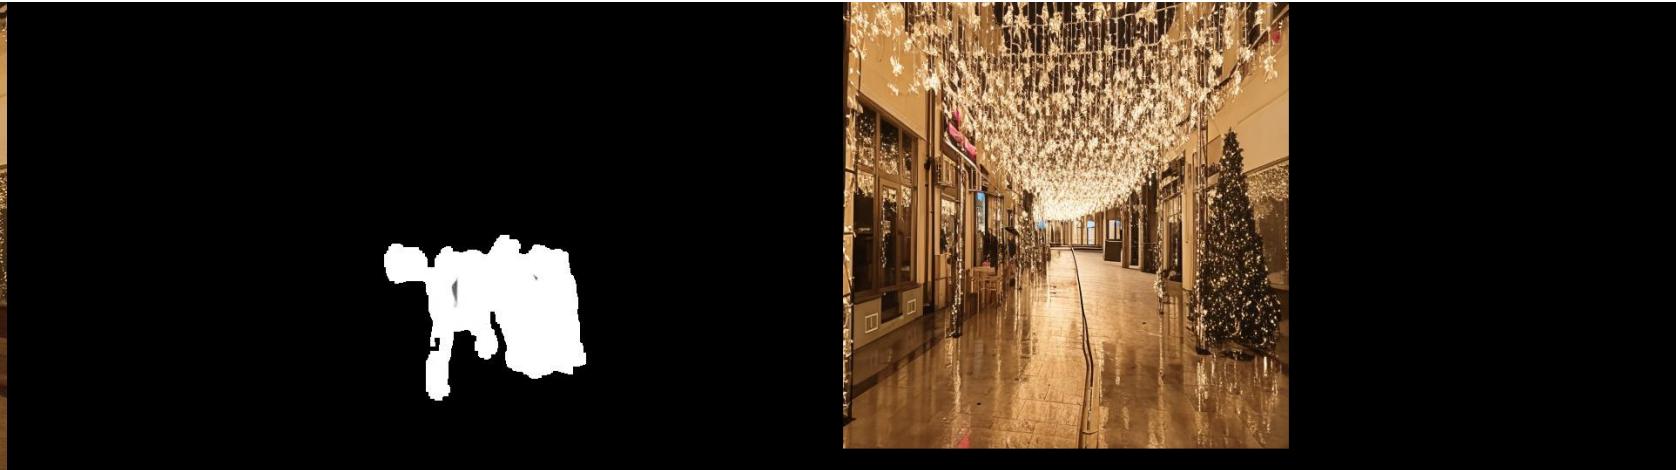
```
from google.colab import drive  
drive.mount('/content/drive')
```

```
import torch  
from diffusers import AutoPipelineForInpainting  
from diffusers.utils import load_image, make_image_grid  
  
pipeline = AutoPipelineForInpainting.from_pretrained(  
    "kandinsky-community/kandinsky-2-2-decoder-inpaint", torch_dtype=torch.float16  
).to("cuda")  
pipeline.enable_model_cpu_offload()  
pipeline.enable_xformers_memory_efficient_attention()
```

```
init_image = load_image("/content/drive/My Drive/Dataset/pics/Craiova.jpg")  
mask_image = load_image("/content/drive/My Drive/Dataset/pics/result_image2.jpg")  
  
prompt = "empty street, with stores and similar floor lighting as the one near 8k"  
negative_prompt = "humans, blue light"  
image = pipeline(prompt=prompt, negative_prompt=negative_prompt, image=init_image, mask_image=mask_image).images[0]  
make_image_grid([init_image, mask_image, image], rows=1, cols=3)
```

Craiova Christmas lights

Initial image – Mask – Inpainting Result



GAN output in paper

Your GAN output



Homework

- Implement a direct convolutional neural network for Fashion MNIST

<https://www.kaggle.com/datasets/zalando-research/fashionmnist>

Or the colored MNIST

<https://www.kaggle.com/datasets/youssifhisham/colored-mnist-dataset/code>

- Implement a pretrained architecture (e.g. ResNet, Inception) for CIFAR-100

<https://keras.io/api/datasets/cifar100/>

- (Optional) Investigate and create a YOLO model for object detection on the data set of your choice
- (Optional) Investigate <https://wandb.ai/> for parametrization and keeping track of experiments





Me

Deep learning
Simple
problem